

# Sheep\_Face\_Classification

April 28, 2024

```
[2]: import numpy as np
import pandas as pd
import os
from PIL import Image
from keras.applications.vgg16 import VGG16, preprocess_input
from tensorflow.keras.applications import VGG19
from keras.preprocessing import image
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Model, load_model
from keras.layers import Dense, GlobalAveragePooling2D, Dropout,
    ↳BatchNormalization, Flatten, Conv2D, Activation, MaxPooling2D
from keras.callbacks import EarlyStopping, Callback, ModelCheckpoint
from sklearn.model_selection import train_test_split, KFold
from tensorflow.keras.utils import to_categorical
import tensorflow as tf
from tensorflow.keras.optimizers import SGD
from sklearn.utils import shuffle
from sklearn.model_selection import StratifiedKFold
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↳f1_score
```

```
[1]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[3]: # Load images and labels
path = '/content/drive/MyDrive/Deep_learning_projects/SheepFaceImages'
sheep_breed_list = os.listdir(path)
print(sheep_breed_list)
```

```
['White Suffolk', 'Poll Dorset', 'Marino', 'Suffolk']
```

```
[4]: def load_and_preprocess_images(path, target_size=(224, 224)):
    img_data_list = []
    labels = []
```

```

for idx, sheep_breed in enumerate(sheep_breed_list):
    sheep_breed_files = os.listdir(os.path.join(path, sheep_breed))
    print(sheep_breed_files)
    for img_file in sheep_breed_files:
        sheep_image_path = os.path.join(path, sheep_breed, img_file)
        try:
            img = image.load_img(sheep_image_path, target_size=target_size)
            img_array = image.img_to_array(img)
            img_array = preprocess_input(img_array)
            img_data_list.append(img_array)
            labels.append(idx)
        except Exception as e:
            print(f"Error loading image {img_file}: {e}")

    # Convert lists to numpy arrays and preprocess
    img_data = np.array(img_data_list)
    labels = np.array(labels)

    # Shuffle data
    img_data, labels = shuffle(img_data, labels, random_state=777)

    return img_data, labels, sheep_breed_list

img_data_path = '/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/
↳sheep_img_data.npy'
labels_path = '/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/
↳sheep_labels.npy'

if os.path.exists(img_data_path) and os.path.exists(labels_path):
    img_data = np.load(img_data_path)
    labels = np.load(labels_path)
else:
    img_data, labels, sheep_breed_list = load_and_preprocess_images(path)
    # Save the processed data
    np.save(img_data_path, img_data)
    np.save(labels_path, labels)

print('Data Shape:', img_data.shape)
print('Labels Shape:', labels.shape)

```

Data Shape: (1680, 224, 224, 3)  
Labels Shape: (1680,)

```

[5]: # One-hot encode labels
labels_categorical = to_categorical(labels, num_classes=len(sheep_breed_list))

print(labels_categorical.shape)

```

(1680, 4)

```
[8]: def build_model():
    base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

    # Freeze first 10 convolutional layers
    for layer in base_model.layers[:15]:
        layer.trainable = False

    x = base_model.output

    # Flatten to prepare for the fully connected layers
    x = Flatten()(x)

    # Fully connected + ReLU
    x = Dense(64, activation='relu')(x)
    x = Dense(64, activation='relu')(x)

    # Adding output layer
    predictions = Dense(len(sheep_breed_list), activation='softmax')(x)

    # Create the final model
    model = Model(inputs=base_model.input, outputs=predictions)

    # Custom learning rates
    lr_mult = {}

    # Set learning rate to 0 for the first 10 layers
    for layer in model.layers[:15]:
        lr_mult[layer.name + '/kernel:0'] = 0.0
        lr_mult[layer.name + '/bias:0'] = 0.0

    # Set learning rate to 0.0001 for the next Conv2D layers and the following layers
    for layer in model.layers[15:-4]:
        if isinstance(layer, Conv2D):
            lr_mult[layer.name + '/kernel:0'] = 0.0001
            lr_mult[layer.name + '/bias:0'] = 0.0001

    # Set learning rate to 10 for the last two dense layers
    for layer in model.layers[-4:]:
        lr_mult[layer.name + '/kernel:0'] = 10
        lr_mult[layer.name + '/bias:0'] = 10

    # Use SGD optimizer with initial learning rate of 1e-4
    optimizer = SGD(learning_rate=1e-4, momentum=0.9)
```

```

    # Compile the model
    model.compile(optimizer=optimizer, loss='categorical_crossentropy',
↪metrics=['accuracy'], loss_weights=[1.] + [10.]*2)

    return model

accuracy_values = []
precision_values = []
recall_values = []
f1_values = []

train_acc_list = []
val_acc_list = []
train_loss_list = []
val_loss_list = []

# KFold Cross-validation
n_splits = 5
kf = KFold(n_splits=n_splits, shuffle=True, random_state=777)
fold_var = 1

for train_index, test_index in kf.split(img_data):
    X_train, X_test = img_data[train_index], img_data[test_index]
    y_train, y_test = labels_categorical[train_index],
↪labels_categorical[test_index]

    model = build_model()

    #early_stop = EarlyStopping(monitor='val_loss', patience=3, verbose=1,
↪mode='min')
    save_path='/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/
↪sheepFace_one.h5'
    checkpoint =
↪ModelCheckpoint(save_path,monitor='val_loss',save_best_only=True,verbose=1)
    history = model.fit(X_train, y_train, batch_size=10, epochs=10, verbose=1,
↪validation_split=0.2, callbacks=[checkpoint])

    print("\n")
    print(f'Evaluating the Test metrics')
    model=load_model(save_path)
    # Evaluate the model
    scores = model.evaluate(X_test, y_test, verbose=1)
    print(f'Test loss for fold {fold_var}: {scores[0]}')
    print(f'Test accuracy for fold {fold_var}: {scores[1]}')

```

```

accuracy_values.append(scores[1]) # Appending accuracy to the list

# Get predictions from the model
y_true = np.argmax(y_test, axis=1) # Convert one-hot encoded labels back
↳to categorical labels
y_pred = np.argmax(model.predict(X_test), axis=1) # Get predictions from
↳the model

# Calculate and append evaluation metrics for this fold
precision = precision_score(y_true, y_pred, average='weighted',
↳zero_division=1)
recall = recall_score(y_true, y_pred, average='weighted')
f1 = f1_score(y_true, y_pred, average='weighted')

precision_values.append(precision)
recall_values.append(recall)
f1_values.append(f1)

# Print evaluation metrics for this fold
print(f'Precision for fold {fold_var}: {precision}')
print(f'Recall for fold {fold_var}: {recall}')
print(f'F1 Score for fold {fold_var}: {f1}')
print("\n")

train_acc_list.append(history.history['accuracy'])
val_acc_list.append(history.history['val_accuracy'])
train_loss_list.append(history.history['loss'])
val_loss_list.append(history.history['val_loss'])

# Increment the fold number
fold_var += 1

# Calculate standard deviation of accuracy across folds
accuracy_std_dev = np.std(accuracy_values)
print(f'Standard deviation of accuracy across folds: {accuracy_std_dev}')

# Calculate average of accuracy, precision, recall, and F1 across folds
avg_accuracy = np.mean(accuracy_values)
avg_precision = np.mean(precision_values)
avg_recall = np.mean(recall_values)
avg_f1 = np.mean(f1_values)

print(f'Average Accuracy across folds: {avg_accuracy}')
print(f'Average Precision across folds: {avg_precision}')
print(f'Average Recall across folds: {avg_recall}')
print(f'Average F1 Score across folds: {avg_f1}')

```

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58889256/58889256 [=====] - 3s 0us/step
Epoch 1/10
108/108 [=====] - ETA: 0s - loss: 1.0960 - accuracy:
0.6521
Epoch 1: val_loss improved from inf to 0.55297, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:
UserWarning: You are saving your model as an HDF5 file via `model.save()`. This
file format is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')`.
    saving_api.save_model(
108/108 [=====] - 31s 227ms/step - loss: 1.0960 -
accuracy: 0.6521 - val_loss: 0.5530 - val_accuracy: 0.7695
Epoch 2/10
107/108 [=====>.] - ETA: 0s - loss: 0.1838 - accuracy:
0.9383
Epoch 2: val_loss improved from 0.55297 to 0.21772, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 8s 75ms/step - loss: 0.1830 -
accuracy: 0.9386 - val_loss: 0.2177 - val_accuracy: 0.9368
Epoch 3/10
107/108 [=====>.] - ETA: 0s - loss: 0.0263 - accuracy:
0.9953
Epoch 3: val_loss improved from 0.21772 to 0.18921, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 8s 75ms/step - loss: 0.0265 -
accuracy: 0.9953 - val_loss: 0.1892 - val_accuracy: 0.9331
Epoch 4/10
107/108 [=====>.] - ETA: 0s - loss: 0.0057 - accuracy:
0.9991
Epoch 4: val_loss improved from 0.18921 to 0.16608, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 8s 76ms/step - loss: 0.0057 -
accuracy: 0.9991 - val_loss: 0.1661 - val_accuracy: 0.9517
Epoch 5/10
107/108 [=====>.] - ETA: 0s - loss: 0.0023 - accuracy:
1.0000
Epoch 5: val_loss did not improve from 0.16608
108/108 [=====] - 8s 73ms/step - loss: 0.0023 -
accuracy: 1.0000 - val_loss: 0.1664 - val_accuracy: 0.9517
Epoch 6/10
107/108 [=====>.] - ETA: 0s - loss: 0.0015 - accuracy:
1.0000
Epoch 6: val_loss improved from 0.16608 to 0.16448, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5

```

```

108/108 [=====] - 9s 86ms/step - loss: 0.0015 -
accuracy: 1.0000 - val_loss: 0.1645 - val_accuracy: 0.9517
Epoch 7/10
107/108 [=====>.] - ETA: 0s - loss: 0.0012 - accuracy:
1.0000
Epoch 7: val_loss improved from 0.16448 to 0.16355, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 8s 76ms/step - loss: 0.0012 -
accuracy: 1.0000 - val_loss: 0.1636 - val_accuracy: 0.9517
Epoch 8/10
107/108 [=====>.] - ETA: 0s - loss: 9.6354e-04 -
accuracy: 1.0000
Epoch 8: val_loss improved from 0.16355 to 0.16288, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 8s 76ms/step - loss: 9.5954e-04 -
accuracy: 1.0000 - val_loss: 0.1629 - val_accuracy: 0.9517
Epoch 9/10
107/108 [=====>.] - ETA: 0s - loss: 8.0315e-04 -
accuracy: 1.0000
Epoch 9: val_loss improved from 0.16288 to 0.16151, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 8s 76ms/step - loss: 8.1136e-04 -
accuracy: 1.0000 - val_loss: 0.1615 - val_accuracy: 0.9517
Epoch 10/10
107/108 [=====>.] - ETA: 0s - loss: 7.0952e-04 -
accuracy: 1.0000
Epoch 10: val_loss did not improve from 0.16151
108/108 [=====] - 8s 74ms/step - loss: 7.1443e-04 -
accuracy: 1.0000 - val_loss: 0.1626 - val_accuracy: 0.9517

```

#### Evaluating the Test metrics

```

11/11 [=====] - 12s 493ms/step - loss: 0.2275 -
accuracy: 0.9345
Test loss for fold 1: 0.22748687863349915
Test accuracy for fold 1: 0.9345238208770752
11/11 [=====] - 1s 122ms/step
Precision for fold 1: 0.9351425705647779
Recall for fold 1: 0.9345238095238095
F1 Score for fold 1: 0.9344968942079089

```

```

Epoch 1/10
107/108 [=====>.] - ETA: 0s - loss: 1.1425 - accuracy:
0.6645
Epoch 1: val_loss improved from inf to 0.45110, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5

```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:
UserWarning: You are saving your model as an HDF5 file via `model.save()`. This
file format is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')`.
```

```
saving_api.save_model(
```

```
108/108 [=====] - 11s 91ms/step - loss: 1.1376 -
accuracy: 0.6660 - val_loss: 0.4511 - val_accuracy: 0.8104
```

```
Epoch 2/10
```

```
107/108 [=====>.] - ETA: 0s - loss: 0.1745 - accuracy:
0.9402
```

```
Epoch 2: val_loss improved from 0.45110 to 0.31211, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
```

```
108/108 [=====] - 8s 78ms/step - loss: 0.1738 -
accuracy: 0.9405 - val_loss: 0.3121 - val_accuracy: 0.8959
```

```
Epoch 3/10
```

```
107/108 [=====>.] - ETA: 0s - loss: 0.0288 - accuracy:
0.9944
```

```
Epoch 3: val_loss improved from 0.31211 to 0.18634, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
```

```
108/108 [=====] - 8s 79ms/step - loss: 0.0293 -
accuracy: 0.9935 - val_loss: 0.1863 - val_accuracy: 0.9219
```

```
Epoch 4/10
```

```
107/108 [=====>.] - ETA: 0s - loss: 0.0064 - accuracy:
1.0000
```

```
Epoch 4: val_loss improved from 0.18634 to 0.16263, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
```

```
108/108 [=====] - 8s 79ms/step - loss: 0.0064 -
accuracy: 1.0000 - val_loss: 0.1626 - val_accuracy: 0.9368
```

```
Epoch 5/10
```

```
107/108 [=====>.] - ETA: 0s - loss: 0.0023 - accuracy:
1.0000
```

```
Epoch 5: val_loss improved from 0.16263 to 0.15170, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
```

```
108/108 [=====] - 9s 79ms/step - loss: 0.0024 -
accuracy: 1.0000 - val_loss: 0.1517 - val_accuracy: 0.9442
```

```
Epoch 6/10
```

```
107/108 [=====>.] - ETA: 0s - loss: 0.0016 - accuracy:
1.0000
```

```
Epoch 6: val_loss improved from 0.15170 to 0.14723, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
```

```
108/108 [=====] - 9s 79ms/step - loss: 0.0016 -
accuracy: 1.0000 - val_loss: 0.1472 - val_accuracy: 0.9517
```

```
Epoch 7/10
```

```
107/108 [=====>.] - ETA: 0s - loss: 0.0012 - accuracy:
1.0000
```

```
Epoch 7: val_loss improved from 0.14723 to 0.14434, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
```



```

108/108 [=====] - 9s 80ms/step - loss: 0.0012 -
accuracy: 1.0000 - val_loss: 0.1443 - val_accuracy: 0.9517
Epoch 8/10
107/108 [=====>.] - ETA: 0s - loss: 0.0010 - accuracy:
1.0000
Epoch 8: val_loss improved from 0.14434 to 0.14239, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 9s 80ms/step - loss: 0.0010 -
accuracy: 1.0000 - val_loss: 0.1424 - val_accuracy: 0.9480
Epoch 9/10
107/108 [=====>.] - ETA: 0s - loss: 8.4061e-04 -
accuracy: 1.0000
Epoch 9: val_loss improved from 0.14239 to 0.14113, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 9s 80ms/step - loss: 8.4538e-04 -
accuracy: 1.0000 - val_loss: 0.1411 - val_accuracy: 0.9517
Epoch 10/10
107/108 [=====>.] - ETA: 0s - loss: 7.3293e-04 -
accuracy: 1.0000
Epoch 10: val_loss improved from 0.14113 to 0.14062, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 9s 80ms/step - loss: 7.3283e-04 -
accuracy: 1.0000 - val_loss: 0.1406 - val_accuracy: 0.9480

```

#### Evaluating the Test metrics

```

11/11 [=====] - 2s 121ms/step - loss: 0.1686 -
accuracy: 0.9613
Test loss for fold 2: 0.16860640048980713
Test accuracy for fold 2: 0.961309552192688
11/11 [=====] - 1s 122ms/step
Precision for fold 2: 0.963667290886392
Recall for fold 2: 0.9613095238095238
F1 Score for fold 2: 0.9613963940006381

```

```

Epoch 1/10
107/108 [=====>.] - ETA: 0s - loss: 1.2414 - accuracy:
0.6467
Epoch 1: val_loss improved from inf to 0.40683, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:
UserWarning: You are saving your model as an HDF5 file via `model.save()`. This
file format is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')`.
  saving_api.save_model(
108/108 [=====] - 12s 104ms/step - loss: 1.2392 -

```

accuracy: 0.6465 - val\_loss: 0.4068 - val\_accuracy: 0.8178  
Epoch 2/10  
107/108 [=====>.] - ETA: 0s - loss: 0.2131 - accuracy: 0.9206  
Epoch 2: val\_loss improved from 0.40683 to 0.24257, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 80ms/step - loss: 0.2122 - accuracy: 0.9209 - val\_loss: 0.2426 - val\_accuracy: 0.9108  
Epoch 3/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0479 - accuracy: 0.9907  
Epoch 3: val\_loss improved from 0.24257 to 0.16371, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 80ms/step - loss: 0.0477 - accuracy: 0.9907 - val\_loss: 0.1637 - val\_accuracy: 0.9480  
Epoch 4/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0080 - accuracy: 0.9991  
Epoch 4: val\_loss improved from 0.16371 to 0.15574, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 80ms/step - loss: 0.0080 - accuracy: 0.9991 - val\_loss: 0.1557 - val\_accuracy: 0.9405  
Epoch 5/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0035 - accuracy: 1.0000  
Epoch 5: val\_loss improved from 0.15574 to 0.14406, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 81ms/step - loss: 0.0036 - accuracy: 1.0000 - val\_loss: 0.1441 - val\_accuracy: 0.9405  
Epoch 6/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0027 - accuracy: 1.0000  
Epoch 6: val\_loss improved from 0.14406 to 0.13473, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 81ms/step - loss: 0.0027 - accuracy: 1.0000 - val\_loss: 0.1347 - val\_accuracy: 0.9480  
Epoch 7/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0019 - accuracy: 1.0000  
Epoch 7: val\_loss did not improve from 0.13473  
108/108 [=====] - 8s 78ms/step - loss: 0.0019 - accuracy: 1.0000 - val\_loss: 0.1361 - val\_accuracy: 0.9442  
Epoch 8/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0015 - accuracy: 1.0000  
Epoch 8: val\_loss did not improve from 0.13473  
108/108 [=====] - 8s 78ms/step - loss: 0.0015 - accuracy: 1.0000 - val\_loss: 0.1407 - val\_accuracy: 0.9480

Epoch 9/10  
 107/108 [=====>.] - ETA: 0s - loss: 0.0013 - accuracy: 1.0000  
 Epoch 9: val\_loss did not improve from 0.13473  
 108/108 [=====] - 8s 78ms/step - loss: 0.0013 - accuracy: 1.0000 - val\_loss: 0.1400 - val\_accuracy: 0.9442  
 Epoch 10/10  
 107/108 [=====>.] - ETA: 0s - loss: 0.0011 - accuracy: 1.0000  
 Epoch 10: val\_loss did not improve from 0.13473  
 108/108 [=====] - 8s 78ms/step - loss: 0.0011 - accuracy: 1.0000 - val\_loss: 0.1381 - val\_accuracy: 0.9442

#### Evaluating the Test metrics

11/11 [=====] - 2s 120ms/step - loss: 0.2372 - accuracy: 0.9286  
 Test loss for fold 3: 0.23715221881866455  
 Test accuracy for fold 3: 0.9285714030265808  
 11/11 [=====] - 1s 123ms/step  
 Precision for fold 3: 0.9309787305944204  
 Recall for fold 3: 0.9285714285714286  
 F1 Score for fold 3: 0.9288600084970795

Epoch 1/10  
 107/108 [=====>.] - ETA: 0s - loss: 1.0911 - accuracy: 0.6682  
 Epoch 1: val\_loss improved from inf to 0.39910, saving model to  
 /content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
 /usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:  
 UserWarning: You are saving your model as an HDF5 file via `model.save()`. This  
 file format is considered legacy. We recommend using instead the native Keras  
 format, e.g. `model.save('my\_model.keras')`.  
 saving\_api.save\_model(  
 108/108 [=====] - 10s 86ms/step - loss: 1.0881 - accuracy: 0.6688 - val\_loss: 0.3991 - val\_accuracy: 0.8290  
 Epoch 2/10  
 107/108 [=====>.] - ETA: 0s - loss: 0.1584 - accuracy: 0.9523  
 Epoch 2: val\_loss improved from 0.39910 to 0.31895, saving model to  
 /content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
 108/108 [=====] - 9s 80ms/step - loss: 0.1579 - accuracy: 0.9526 - val\_loss: 0.3190 - val\_accuracy: 0.8922  
 Epoch 3/10  
 107/108 [=====>.] - ETA: 0s - loss: 0.0258 - accuracy: 0.9944

Epoch 3: val\_loss improved from 0.31895 to 0.15681, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 80ms/step - loss: 0.0258 -  
accuracy: 0.9944 - val\_loss: 0.1568 - val\_accuracy: 0.9480  
Epoch 4/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0048 - accuracy:  
1.0000  
Epoch 4: val\_loss improved from 0.15681 to 0.14781, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 81ms/step - loss: 0.0048 -  
accuracy: 1.0000 - val\_loss: 0.1478 - val\_accuracy: 0.9480  
Epoch 5/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0025 - accuracy:  
1.0000  
Epoch 5: val\_loss improved from 0.14781 to 0.14488, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 81ms/step - loss: 0.0025 -  
accuracy: 1.0000 - val\_loss: 0.1449 - val\_accuracy: 0.9442  
Epoch 6/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0018 - accuracy:  
1.0000  
Epoch 6: val\_loss improved from 0.14488 to 0.14384, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 81ms/step - loss: 0.0018 -  
accuracy: 1.0000 - val\_loss: 0.1438 - val\_accuracy: 0.9480  
Epoch 7/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0014 - accuracy:  
1.0000  
Epoch 7: val\_loss improved from 0.14384 to 0.14070, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 81ms/step - loss: 0.0014 -  
accuracy: 1.0000 - val\_loss: 0.1407 - val\_accuracy: 0.9517  
Epoch 8/10  
107/108 [=====>.] - ETA: 0s - loss: 0.0012 - accuracy:  
1.0000  
Epoch 8: val\_loss improved from 0.14070 to 0.14065, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 9s 81ms/step - loss: 0.0012 -  
accuracy: 1.0000 - val\_loss: 0.1407 - val\_accuracy: 0.9480  
Epoch 9/10  
107/108 [=====>.] - ETA: 0s - loss: 9.9359e-04 -  
accuracy: 1.0000  
Epoch 9: val\_loss did not improve from 0.14065  
108/108 [=====] - 8s 78ms/step - loss: 9.9759e-04 -  
accuracy: 1.0000 - val\_loss: 0.1429 - val\_accuracy: 0.9480  
Epoch 10/10  
107/108 [=====>.] - ETA: 0s - loss: 8.7069e-04 -  
accuracy: 1.0000

Epoch 10: val\_loss improved from 0.14065 to 0.14020, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFace\_Data/sheepFace\_one.h5  
108/108 [=====] - 10s 90ms/step - loss: 8.6784e-04 -  
accuracy: 1.0000 - val\_loss: 0.1402 - val\_accuracy: 0.9517

```
11/11 [=====] - 2s 125ms/step - loss: 0.1534 - accuracy: 0.9554
```

Epoch 1/10

```

/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 9s 81ms/step - loss: 0.0055 -
accuracy: 1.0000 - val_loss: 0.1560 - val_accuracy: 0.9554
Epoch 5/10
107/108 [=====>.] - ETA: 0s - loss: 0.0030 - accuracy:
1.0000
Epoch 5: val_loss improved from 0.15597 to 0.14799, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/sheepFace_one.h5
108/108 [=====] - 9s 82ms/step - loss: 0.0030 -
accuracy: 1.0000 - val_loss: 0.1480 - val_accuracy: 0.9628
Epoch 6/10
107/108 [=====>.] - ETA: 0s - loss: 0.0019 - accuracy:
1.0000
Epoch 6: val_loss did not improve from 0.14799
108/108 [=====] - 8s 79ms/step - loss: 0.0019 -
accuracy: 1.0000 - val_loss: 0.1517 - val_accuracy: 0.9591
Epoch 7/10
107/108 [=====>.] - ETA: 0s - loss: 0.0014 - accuracy:
1.0000
Epoch 7: val_loss did not improve from 0.14799
108/108 [=====] - 8s 79ms/step - loss: 0.0014 -
accuracy: 1.0000 - val_loss: 0.1517 - val_accuracy: 0.9591
Epoch 8/10
107/108 [=====>.] - ETA: 0s - loss: 0.0012 - accuracy:
1.0000
Epoch 8: val_loss did not improve from 0.14799
108/108 [=====] - 8s 79ms/step - loss: 0.0012 -
accuracy: 1.0000 - val_loss: 0.1511 - val_accuracy: 0.9628
Epoch 9/10
107/108 [=====>.] - ETA: 0s - loss: 9.7564e-04 -
accuracy: 1.0000
Epoch 9: val_loss did not improve from 0.14799
108/108 [=====] - 9s 79ms/step - loss: 9.7703e-04 -
accuracy: 1.0000 - val_loss: 0.1539 - val_accuracy: 0.9591
Epoch 10/10
107/108 [=====>.] - ETA: 0s - loss: 8.4613e-04 -
accuracy: 1.0000
Epoch 10: val_loss did not improve from 0.14799
108/108 [=====] - 9s 79ms/step - loss: 8.4295e-04 -
accuracy: 1.0000 - val_loss: 0.1547 - val_accuracy: 0.9628

```

Evaluating the Test metrics

```

11/11 [=====] - 2s 124ms/step - loss: 0.2970 -
accuracy: 0.9077
Test loss for fold 5: 0.2970311641693115
Test accuracy for fold 5: 0.9077380895614624
11/11 [=====] - 1s 124ms/step

```

Precision for fold 5: 0.908393509773728  
Recall for fold 5: 0.9077380952380952  
F1 Score for fold 5: 0.9059085571465838

Standard deviation of accuracy across folds: 0.01928792794246443  
Average Accuracy across folds: 0.9375  
Average Precision across folds: 0.9387839152592576  
Average Recall across folds: 0.9375  
Average F1 Score across folds: 0.937218991490583

```
[9]: import matplotlib.pyplot as plt

# Determine the maximum length among all lists
max_length = max(len(train_acc) for train_acc in train_acc_list)

# Pad the shorter lists with zeros to match the maximum length
for i in range(n_splits):
    train_acc_list[i] += [0] * (max_length - len(train_acc_list[i]))
    val_acc_list[i] += [0] * (max_length - len(val_acc_list[i]))
    train_loss_list[i] += [0] * (max_length - len(train_loss_list[i]))
    val_loss_list[i] += [0] * (max_length - len(val_loss_list[i]))

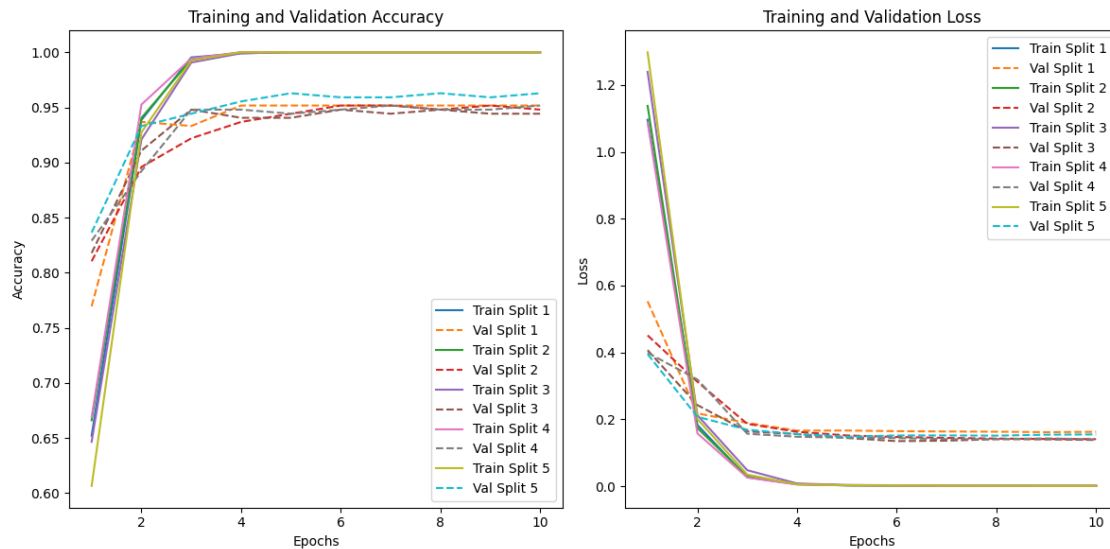
# Plotting
epochs = range(1, max_length + 1)

plt.figure(figsize=(12, 6))

# Plotting Training and Validation Accuracy
plt.subplot(1, 2, 1)
for i in range(n_splits):
    plt.plot(epochs, train_acc_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_acc_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Plotting Training and Validation Loss
plt.subplot(1, 2, 2)
for i in range(n_splits):
    plt.plot(epochs, train_loss_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_loss_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
```

```
plt.tight_layout()
plt.show()
```



```
[15]: def build_model():
    base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

    # Freeze first 10 convolutional layers
    for layer in base_model.layers[:15]:
        layer.trainable = False

    x = base_model.output

    # Adding Conv2D + ReLU Layers
    x = Conv2D(512, (3, 3), padding='same')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = Conv2D(512, (3, 3), padding='same', activation='relu')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = Conv2D(512, (3, 3), padding='same', activation='relu')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)

    # Adding Max Pooling
    x = MaxPooling2D((2, 2), strides=(2, 2))(x)
```



```

# Add GlobalAveragePooling2D layer
x = GlobalAveragePooling2D()(x)

# Replacing the last three layers with new fully connected layers
x = Dense(256)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)

x = Dense(256)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)

predictions = Dense(len(sheep_breed_list), activation='softmax')(x) # Add
↪output layer

model_normalizeddd = Model(inputs=base_model.input, outputs=predictions)

# Custom learning rates
lr_mult = {}

# Set learning rate to 0 for the first 10 layers
for layer in model_normalizeddd.layers[:15]:
    lr_mult[layer.name + '/kernel:0'] = 0.0
    lr_mult[layer.name + '/bias:0'] = 0.0

# Set learning rate to 0.0001 for the next Conv2D layers and the following
↪layers
for layer in model_normalizeddd.layers[15:30]:
    if isinstance(layer, Conv2D):
        lr_mult[layer.name + '/kernel:0'] = 0.0001
        lr_mult[layer.name + '/bias:0'] = 0.0001

# Set learning rate to 10 for the last two dense layers
for layer in model_normalizeddd.layers[30:]:
    lr_mult[layer.name + '/kernel:0'] = 1
    lr_mult[layer.name + '/bias:0'] = 1

# Use SGD optimizer with initial learning rate of 1e-4
optimizer = SGD(learning_rate=1e-4, momentum=0.9)

# Compile the model
model_normalizeddd.compile(optimizer=optimizer,
↪loss='categorical_crossentropy', metrics=['accuracy'], loss_weights=[1.] +
↪[10.]*2)

```

```

    return model_normalizeddd

accuracy_values = []
precision_values = []
recall_values = []
f1_values = []

train_acc_list = []
val_acc_list = []
train_loss_list = []
val_loss_list = []

# KFold Cross-validation
n_splits = 5
kf = KFold(n_splits=n_splits, shuffle=True, random_state=777)
fold_var = 1

for train_index, test_index in kf.split(img_data):
    X_train, X_test = img_data[train_index], img_data[test_index]
    y_train, y_test = labels_categorical[train_index],
    ↪ labels_categorical[test_index]

    model_normalizeddd = build_model()

    # Define the ImageDataGenerator for data augmentation
    datagen = ImageDataGenerator(
        width_shift_range=0.1, # translate horizontally by 10% of total width
        height_shift_range=0.1, # translate vertically by 10% of total height
        fill_mode='nearest' # strategy for filling in newly created pixels
    )

    # Generate augmented images batches during training
    train_datagen = datagen.flow(X_train, y_train, batch_size=10)

    #early_stop = EarlyStopping(monitor='val_loss', patience=3, verbose=1, ↪
    ↪ mode='min')
    save_path='/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/
    ↪ sheepFace_four.keras'
    checkpoint = ↪
    ↪ ModelCheckpoint(save_path,monitor='val_loss',save_best_only=True,verbose=0)
    history = model_normalizeddd.fit(train_datagen, epochs=20, ↪
    ↪ validation_data=(X_test,y_test), verbose=1, callbacks=[checkpoint])

    print("\n")
    print(f'Evaluating the Test metrics')
    model_normalizeddd=load_model(save_path)
    # Evaluate the model

```

```

scores = model_normalizedd.evaluate(X_test, y_test, verbose=1)
print(f'Test loss for fold {fold_var}: {scores[0]}')
print(f'Test accuracy for fold {fold_var}: {scores[1]}')

accuracy_values.append(scores[1]) # Appending accuracy to the list

# Get predictions from the model
y_true = np.argmax(y_test, axis=1) # Convert one-hot encoded labels back
↳to categorical labels
y_pred = np.argmax(model_normalizedd.predict(X_test), axis=1) # Get
↳predictions from the model

# Calculate and append evaluation metrics for this fold
precision = precision_score(y_true, y_pred, average='weighted',
↳zero_division=1)
recall = recall_score(y_true, y_pred, average='weighted')
f1 = f1_score(y_true, y_pred, average='weighted')

precision_values.append(precision)
recall_values.append(recall)
f1_values.append(f1)

# Print evaluation metrics for this fold
print(f'Precision for fold {fold_var}: {precision}')
print(f'Recall for fold {fold_var}: {recall}')
print(f'F1 Score for fold {fold_var}: {f1}')

train_acc_list.append(history.history['accuracy'])
val_acc_list.append(history.history['val_accuracy'])
train_loss_list.append(history.history['loss'])
val_loss_list.append(history.history['val_loss'])
print("\n")

# Increment the fold number
fold_var += 1

# Calculate standard deviation of accuracy across folds
accuracy_std_dev = np.std(accuracy_values)
print(f'Standard deviation of accuracy across folds: {accuracy_std_dev}')

# Calculate average of accuracy, precision, recall, and F1 across folds
avg_accuracy = np.mean(accuracy_values)
avg_precision = np.mean(precision_values)
avg_recall = np.mean(recall_values)
avg_f1 = np.mean(f1_values)

print(f'Average Accuracy across folds: {avg_accuracy}')

```

```
print(f'Average Precision across folds: {avg_precision}')
```

```
print(f'Average Recall across folds: {avg_recall}')
```

```
print(f'Average F1 Score across folds: {avg_f1}')
```

Epoch 1/20

135/135 [=====] - 18s 112ms/step - loss: 1.4963 - accuracy: 0.3884 - val\_loss: 1.0836 - val\_accuracy: 0.5327

Epoch 2/20

135/135 [=====] - 15s 110ms/step - loss: 1.0848 - accuracy: 0.5603 - val\_loss: 0.6992 - val\_accuracy: 0.7530

Epoch 3/20

135/135 [=====] - 14s 102ms/step - loss: 0.8431 - accuracy: 0.6704 - val\_loss: 0.9500 - val\_accuracy: 0.6071

Epoch 4/20

135/135 [=====] - 15s 109ms/step - loss: 0.6366 - accuracy: 0.7552 - val\_loss: 0.4003 - val\_accuracy: 0.8512

Epoch 5/20

135/135 [=====] - 15s 109ms/step - loss: 0.5589 - accuracy: 0.7917 - val\_loss: 0.3772 - val\_accuracy: 0.8571

Epoch 6/20

135/135 [=====] - 14s 106ms/step - loss: 0.4512 - accuracy: 0.8519 - val\_loss: 0.5725 - val\_accuracy: 0.8006

Epoch 7/20

135/135 [=====] - 15s 109ms/step - loss: 0.4128 - accuracy: 0.8504 - val\_loss: 0.2462 - val\_accuracy: 0.9196

Epoch 8/20

135/135 [=====] - 14s 103ms/step - loss: 0.3677 - accuracy: 0.8810 - val\_loss: 0.2631 - val\_accuracy: 0.9226

Epoch 9/20

135/135 [=====] - 15s 110ms/step - loss: 0.3378 - accuracy: 0.8839 - val\_loss: 0.2285 - val\_accuracy: 0.9405

Epoch 10/20

135/135 [=====] - 15s 109ms/step - loss: 0.3156 - accuracy: 0.8996 - val\_loss: 0.2204 - val\_accuracy: 0.9315

Epoch 11/20

135/135 [=====] - 14s 106ms/step - loss: 0.2666 - accuracy: 0.9196 - val\_loss: 0.2230 - val\_accuracy: 0.9167

Epoch 12/20

135/135 [=====] - 15s 109ms/step - loss: 0.2717 - accuracy: 0.9122 - val\_loss: 0.1781 - val\_accuracy: 0.9554

Epoch 13/20

135/135 [=====] - 14s 106ms/step - loss: 0.2107 - accuracy: 0.9368 - val\_loss: 0.1963 - val\_accuracy: 0.9375

Epoch 14/20

135/135 [=====] - 15s 108ms/step - loss: 0.2298 - accuracy: 0.9278 - val\_loss: 0.1557 - val\_accuracy: 0.9673

Epoch 15/20

135/135 [=====] - 14s 105ms/step - loss: 0.2024 -  
accuracy: 0.9345 - val\_loss: 0.1844 - val\_accuracy: 0.9464  
Epoch 16/20  
135/135 [=====] - 15s 110ms/step - loss: 0.1863 -  
accuracy: 0.9405 - val\_loss: 0.1517 - val\_accuracy: 0.9613  
Epoch 17/20  
135/135 [=====] - 15s 110ms/step - loss: 0.1719 -  
accuracy: 0.9464 - val\_loss: 0.1282 - val\_accuracy: 0.9762  
Epoch 18/20  
135/135 [=====] - 14s 104ms/step - loss: 0.1531 -  
accuracy: 0.9539 - val\_loss: 0.1640 - val\_accuracy: 0.9494  
Epoch 19/20  
135/135 [=====] - 14s 105ms/step - loss: 0.1667 -  
accuracy: 0.9568 - val\_loss: 0.1479 - val\_accuracy: 0.9702  
Epoch 20/20  
135/135 [=====] - 14s 106ms/step - loss: 0.1302 -  
accuracy: 0.9680 - val\_loss: 0.1393 - val\_accuracy: 0.9673

Evaluating the Test metrics

11/11 [=====] - 2s 125ms/step - loss: 0.1282 -  
accuracy: 0.9762  
Test loss for fold 1: 0.12819179892539978  
Test accuracy for fold 1: 0.976190447807312  
11/11 [=====] - 1s 126ms/step  
Precision for fold 1: 0.9763740521090573  
Recall for fold 1: 0.9761904761904762  
F1 Score for fold 1: 0.9761648922458949

Epoch 1/20  
135/135 [=====] - 17s 112ms/step - loss: 1.4216 -  
accuracy: 0.4018 - val\_loss: 1.0859 - val\_accuracy: 0.5714  
Epoch 2/20  
135/135 [=====] - 15s 110ms/step - loss: 0.9931 -  
accuracy: 0.6004 - val\_loss: 0.6225 - val\_accuracy: 0.7113  
Epoch 3/20  
135/135 [=====] - 15s 110ms/step - loss: 0.7850 -  
accuracy: 0.6868 - val\_loss: 0.4340 - val\_accuracy: 0.8244  
Epoch 4/20  
135/135 [=====] - 15s 111ms/step - loss: 0.5969 -  
accuracy: 0.7812 - val\_loss: 0.3238 - val\_accuracy: 0.8988  
Epoch 5/20  
135/135 [=====] - 15s 110ms/step - loss: 0.5076 -  
accuracy: 0.8058 - val\_loss: 0.2617 - val\_accuracy: 0.9196  
Epoch 6/20  
135/135 [=====] - 14s 105ms/step - loss: 0.4539 -  
accuracy: 0.8296 - val\_loss: 0.2634 - val\_accuracy: 0.9107

Epoch 7/20  
135/135 [=====] - 15s 108ms/step - loss: 0.3958 - accuracy: 0.8586 - val\_loss: 0.1421 - val\_accuracy: 0.9673  
Epoch 8/20  
135/135 [=====] - 14s 105ms/step - loss: 0.3693 - accuracy: 0.8847 - val\_loss: 0.1731 - val\_accuracy: 0.9464  
Epoch 9/20  
135/135 [=====] - 14s 104ms/step - loss: 0.3281 - accuracy: 0.8906 - val\_loss: 0.1751 - val\_accuracy: 0.9405  
Epoch 10/20  
135/135 [=====] - 15s 108ms/step - loss: 0.2552 - accuracy: 0.9271 - val\_loss: 0.1268 - val\_accuracy: 0.9554  
Epoch 11/20  
135/135 [=====] - 14s 105ms/step - loss: 0.2250 - accuracy: 0.9286 - val\_loss: 0.1692 - val\_accuracy: 0.9345  
Epoch 12/20  
135/135 [=====] - 15s 109ms/step - loss: 0.2719 - accuracy: 0.9159 - val\_loss: 0.1008 - val\_accuracy: 0.9702  
Epoch 13/20  
135/135 [=====] - 14s 103ms/step - loss: 0.2054 - accuracy: 0.9427 - val\_loss: 0.1370 - val\_accuracy: 0.9494  
Epoch 14/20  
135/135 [=====] - 14s 105ms/step - loss: 0.2200 - accuracy: 0.9345 - val\_loss: 0.1213 - val\_accuracy: 0.9583  
Epoch 15/20  
135/135 [=====] - 14s 105ms/step - loss: 0.1942 - accuracy: 0.9427 - val\_loss: 0.1017 - val\_accuracy: 0.9702  
Epoch 16/20  
135/135 [=====] - 15s 108ms/step - loss: 0.1608 - accuracy: 0.9516 - val\_loss: 0.0799 - val\_accuracy: 0.9821  
Epoch 17/20  
135/135 [=====] - 14s 106ms/step - loss: 0.1830 - accuracy: 0.9472 - val\_loss: 0.1207 - val\_accuracy: 0.9554  
Epoch 18/20  
135/135 [=====] - 14s 102ms/step - loss: 0.1508 - accuracy: 0.9546 - val\_loss: 0.1240 - val\_accuracy: 0.9613  
Epoch 19/20  
135/135 [=====] - 15s 109ms/step - loss: 0.1591 - accuracy: 0.9554 - val\_loss: 0.0734 - val\_accuracy: 0.9762  
Epoch 20/20  
135/135 [=====] - 14s 105ms/step - loss: 0.1386 - accuracy: 0.9665 - val\_loss: 0.1089 - val\_accuracy: 0.9524

Evaluating the Test metrics

11/11 [=====] - 2s 128ms/step - loss: 0.0734 - accuracy: 0.9762

Test loss for fold 2: 0.07339658588171005

Test accuracy for fold 2: 0.976190447807312  
11/11 [=====] - 2s 130ms/step  
Precision for fold 2: 0.9763802209676707  
Recall for fold 2: 0.9761904761904762  
F1 Score for fold 2: 0.9761993841087628

Epoch 1/20  
135/135 [=====] - 18s 114ms/step - loss: 1.5529 -  
accuracy: 0.3787 - val\_loss: 1.2174 - val\_accuracy: 0.3423  
Epoch 2/20  
135/135 [=====] - 15s 111ms/step - loss: 1.0124 -  
accuracy: 0.5863 - val\_loss: 0.8827 - val\_accuracy: 0.5833  
Epoch 3/20  
135/135 [=====] - 14s 104ms/step - loss: 0.7682 -  
accuracy: 0.6994 - val\_loss: 0.9962 - val\_accuracy: 0.5863  
Epoch 4/20  
135/135 [=====] - 15s 108ms/step - loss: 0.6080 -  
accuracy: 0.7634 - val\_loss: 0.4824 - val\_accuracy: 0.8125  
Epoch 5/20  
135/135 [=====] - 14s 104ms/step - loss: 0.5265 -  
accuracy: 0.8207 - val\_loss: 0.5382 - val\_accuracy: 0.7738  
Epoch 6/20  
135/135 [=====] - 15s 109ms/step - loss: 0.4607 -  
accuracy: 0.8363 - val\_loss: 0.2290 - val\_accuracy: 0.9226  
Epoch 7/20  
135/135 [=====] - 14s 103ms/step - loss: 0.3941 -  
accuracy: 0.8668 - val\_loss: 0.2806 - val\_accuracy: 0.9107  
Epoch 8/20  
135/135 [=====] - 14s 104ms/step - loss: 0.3529 -  
accuracy: 0.8772 - val\_loss: 0.2308 - val\_accuracy: 0.9315  
Epoch 9/20  
135/135 [=====] - 15s 108ms/step - loss: 0.2866 -  
accuracy: 0.8996 - val\_loss: 0.1788 - val\_accuracy: 0.9435  
Epoch 10/20  
135/135 [=====] - 15s 109ms/step - loss: 0.2476 -  
accuracy: 0.9204 - val\_loss: 0.1663 - val\_accuracy: 0.9583  
Epoch 11/20  
135/135 [=====] - 14s 105ms/step - loss: 0.2449 -  
accuracy: 0.9286 - val\_loss: 0.1998 - val\_accuracy: 0.9464  
Epoch 12/20  
135/135 [=====] - 14s 104ms/step - loss: 0.2849 -  
accuracy: 0.9018 - val\_loss: 0.2388 - val\_accuracy: 0.9256  
Epoch 13/20  
135/135 [=====] - 14s 104ms/step - loss: 0.2358 -  
accuracy: 0.9204 - val\_loss: 0.1670 - val\_accuracy: 0.9494  
Epoch 14/20  
135/135 [=====] - 14s 102ms/step - loss: 0.1901 -

accuracy: 0.9457 - val\_loss: 0.1879 - val\_accuracy: 0.9464  
Epoch 15/20  
135/135 [=====] - 15s 109ms/step - loss: 0.1850 -  
accuracy: 0.9464 - val\_loss: 0.1614 - val\_accuracy: 0.9583  
Epoch 16/20  
135/135 [=====] - 15s 110ms/step - loss: 0.1727 -  
accuracy: 0.9516 - val\_loss: 0.1476 - val\_accuracy: 0.9613  
Epoch 17/20  
135/135 [=====] - 14s 102ms/step - loss: 0.1994 -  
accuracy: 0.9382 - val\_loss: 0.1584 - val\_accuracy: 0.9673  
Epoch 18/20  
135/135 [=====] - 14s 105ms/step - loss: 0.1652 -  
accuracy: 0.9539 - val\_loss: 0.1660 - val\_accuracy: 0.9673  
Epoch 19/20  
135/135 [=====] - 14s 104ms/step - loss: 0.1677 -  
accuracy: 0.9479 - val\_loss: 0.2092 - val\_accuracy: 0.9494  
Epoch 20/20  
135/135 [=====] - 14s 105ms/step - loss: 0.1446 -  
accuracy: 0.9561 - val\_loss: 0.2040 - val\_accuracy: 0.9345

#### Evaluating the Test metrics

11/11 [=====] - 2s 127ms/step - loss: 0.1476 -  
accuracy: 0.9613  
Test loss for fold 3: 0.1475822627544403  
Test accuracy for fold 3: 0.961309552192688  
11/11 [=====] - 2s 132ms/step  
Precision for fold 3: 0.9615065524279209  
Recall for fold 3: 0.9613095238095238  
F1 Score for fold 3: 0.9612206995201181

Epoch 1/20  
135/135 [=====] - 18s 114ms/step - loss: 1.4243 -  
accuracy: 0.3921 - val\_loss: 0.9675 - val\_accuracy: 0.6845  
Epoch 2/20  
135/135 [=====] - 15s 109ms/step - loss: 0.9540 -  
accuracy: 0.6071 - val\_loss: 0.6019 - val\_accuracy: 0.8065  
Epoch 3/20  
135/135 [=====] - 15s 109ms/step - loss: 0.8068 -  
accuracy: 0.6786 - val\_loss: 0.3997 - val\_accuracy: 0.8661  
Epoch 4/20  
135/135 [=====] - 15s 108ms/step - loss: 0.6091 -  
accuracy: 0.7723 - val\_loss: 0.2613 - val\_accuracy: 0.9196  
Epoch 5/20  
135/135 [=====] - 15s 110ms/step - loss: 0.4858 -  
accuracy: 0.8140 - val\_loss: 0.2122 - val\_accuracy: 0.9494  
Epoch 6/20



135/135 [=====] - 14s 103ms/step - loss: 0.4362 - accuracy: 0.8415 - val\_loss: 0.2130 - val\_accuracy: 0.9196  
Epoch 7/20  
135/135 [=====] - 15s 109ms/step - loss: 0.3833 - accuracy: 0.8616 - val\_loss: 0.1938 - val\_accuracy: 0.9464  
Epoch 8/20  
135/135 [=====] - 15s 108ms/step - loss: 0.3593 - accuracy: 0.8780 - val\_loss: 0.1625 - val\_accuracy: 0.9583  
Epoch 9/20  
135/135 [=====] - 15s 109ms/step - loss: 0.3054 - accuracy: 0.8973 - val\_loss: 0.1072 - val\_accuracy: 0.9732  
Epoch 10/20  
135/135 [=====] - 15s 110ms/step - loss: 0.2759 - accuracy: 0.9174 - val\_loss: 0.0928 - val\_accuracy: 0.9851  
Epoch 11/20  
135/135 [=====] - 15s 109ms/step - loss: 0.2609 - accuracy: 0.9137 - val\_loss: 0.0861 - val\_accuracy: 0.9881  
Epoch 12/20  
135/135 [=====] - 15s 109ms/step - loss: 0.2854 - accuracy: 0.9129 - val\_loss: 0.0818 - val\_accuracy: 0.9851  
Epoch 13/20  
135/135 [=====] - 15s 110ms/step - loss: 0.2360 - accuracy: 0.9353 - val\_loss: 0.0762 - val\_accuracy: 0.9970  
Epoch 14/20  
135/135 [=====] - 14s 103ms/step - loss: 0.2162 - accuracy: 0.9345 - val\_loss: 0.1105 - val\_accuracy: 0.9643  
Epoch 15/20  
135/135 [=====] - 14s 104ms/step - loss: 0.1932 - accuracy: 0.9427 - val\_loss: 0.0934 - val\_accuracy: 0.9762  
Epoch 16/20  
135/135 [=====] - 14s 104ms/step - loss: 0.1851 - accuracy: 0.9420 - val\_loss: 0.0854 - val\_accuracy: 0.9792  
Epoch 17/20  
135/135 [=====] - 15s 108ms/step - loss: 0.1618 - accuracy: 0.9546 - val\_loss: 0.0583 - val\_accuracy: 0.9851  
Epoch 18/20  
135/135 [=====] - 15s 109ms/step - loss: 0.1531 - accuracy: 0.9568 - val\_loss: 0.0548 - val\_accuracy: 0.9881  
Epoch 19/20  
135/135 [=====] - 14s 105ms/step - loss: 0.1699 - accuracy: 0.9524 - val\_loss: 0.1322 - val\_accuracy: 0.9464  
Epoch 20/20  
135/135 [=====] - 15s 108ms/step - loss: 0.1313 - accuracy: 0.9650 - val\_loss: 0.0408 - val\_accuracy: 0.9881

Evaluating the Test metrics

11/11 [=====] - 2s 129ms/step - loss: 0.0408 -

accuracy: 0.9881  
Test loss for fold 4: 0.04082081839442253  
Test accuracy for fold 4: 0.988095223903656  
11/11 [=====] - 2s 133ms/step  
Precision for fold 4: 0.9881338899196043  
Recall for fold 4: 0.9880952380952381  
F1 Score for fold 4: 0.9880961589106386

Epoch 1/20  
135/135 [=====] - 18s 114ms/step - loss: 1.4524 -  
accuracy: 0.4211 - val\_loss: 1.0209 - val\_accuracy: 0.5685  
Epoch 2/20  
135/135 [=====] - 15s 108ms/step - loss: 0.9361 -  
accuracy: 0.6124 - val\_loss: 0.8006 - val\_accuracy: 0.6518  
Epoch 3/20  
135/135 [=====] - 15s 110ms/step - loss: 0.6725 -  
accuracy: 0.7433 - val\_loss: 0.4505 - val\_accuracy: 0.8423  
Epoch 4/20  
135/135 [=====] - 15s 109ms/step - loss: 0.5771 -  
accuracy: 0.7812 - val\_loss: 0.3741 - val\_accuracy: 0.8869  
Epoch 5/20  
135/135 [=====] - 15s 108ms/step - loss: 0.4611 -  
accuracy: 0.8400 - val\_loss: 0.3246 - val\_accuracy: 0.8958  
Epoch 6/20  
135/135 [=====] - 15s 108ms/step - loss: 0.4111 -  
accuracy: 0.8519 - val\_loss: 0.2083 - val\_accuracy: 0.9435  
Epoch 7/20  
135/135 [=====] - 15s 109ms/step - loss: 0.3787 -  
accuracy: 0.8728 - val\_loss: 0.1946 - val\_accuracy: 0.9554  
Epoch 8/20  
135/135 [=====] - 15s 109ms/step - loss: 0.3230 -  
accuracy: 0.8943 - val\_loss: 0.1591 - val\_accuracy: 0.9732  
Epoch 9/20  
135/135 [=====] - 14s 104ms/step - loss: 0.2851 -  
accuracy: 0.9055 - val\_loss: 0.1967 - val\_accuracy: 0.9345  
Epoch 10/20  
135/135 [=====] - 15s 109ms/step - loss: 0.2763 -  
accuracy: 0.9085 - val\_loss: 0.1501 - val\_accuracy: 0.9732  
Epoch 11/20  
135/135 [=====] - 15s 110ms/step - loss: 0.2476 -  
accuracy: 0.9226 - val\_loss: 0.1378 - val\_accuracy: 0.9702  
Epoch 12/20  
135/135 [=====] - 15s 109ms/step - loss: 0.2606 -  
accuracy: 0.9174 - val\_loss: 0.1322 - val\_accuracy: 0.9762  
Epoch 13/20  
135/135 [=====] - 16s 117ms/step - loss: 0.2354 -  
accuracy: 0.9219 - val\_loss: 0.1143 - val\_accuracy: 0.9821

```

Epoch 14/20
135/135 [=====] - 16s 116ms/step - loss: 0.2086 -
accuracy: 0.9345 - val_loss: 0.1184 - val_accuracy: 0.9762
Epoch 15/20
135/135 [=====] - 16s 116ms/step - loss: 0.1916 -
accuracy: 0.9427 - val_loss: 0.0975 - val_accuracy: 0.9821
Epoch 16/20
135/135 [=====] - 15s 111ms/step - loss: 0.1782 -
accuracy: 0.9501 - val_loss: 0.0946 - val_accuracy: 0.9821
Epoch 17/20
135/135 [=====] - 14s 104ms/step - loss: 0.2007 -
accuracy: 0.9397 - val_loss: 0.2015 - val_accuracy: 0.9256
Epoch 18/20
135/135 [=====] - 14s 105ms/step - loss: 0.1738 -
accuracy: 0.9457 - val_loss: 0.1833 - val_accuracy: 0.9405
Epoch 19/20
135/135 [=====] - 14s 106ms/step - loss: 0.1581 -
accuracy: 0.9576 - val_loss: 0.1255 - val_accuracy: 0.9702
Epoch 20/20
135/135 [=====] - 15s 108ms/step - loss: 0.1555 -
accuracy: 0.9546 - val_loss: 0.0929 - val_accuracy: 0.9821

```

Evaluating the Test metrics

```

11/11 [=====] - 2s 128ms/step - loss: 0.0929 -
accuracy: 0.9821
Test loss for fold 5: 0.09289499372243881
Test accuracy for fold 5: 0.9821428656578064
11/11 [=====] - 2s 130ms/step
Precision for fold 5: 0.9822686116700202
Recall for fold 5: 0.9821428571428571
F1 Score for fold 5: 0.9821862673879936

```

```

Standard deviation of accuracy across folds: 0.00890869638175907
Average Accuracy across folds: 0.9767857074737549
Average Precision across folds: 0.9769326654188546
Average Recall across folds: 0.9767857142857143
Average F1 Score across folds: 0.9767734804346816

```

```

[16]: import matplotlib.pyplot as plt

# Determine the maximum length among all lists
max_length = max(len(train_acc) for train_acc in train_acc_list)

# Pad the shorter lists with zeros to match the maximum length
for i in range(n_splits):

```

```

train_acc_list[i] += [0] * (max_length - len(train_acc_list[i]))
val_acc_list[i] += [0] * (max_length - len(val_acc_list[i]))
train_loss_list[i] += [0] * (max_length - len(train_loss_list[i]))
val_loss_list[i] += [0] * (max_length - len(val_loss_list[i]))

# Plotting
epochs = range(1, max_length + 1)


plt.figure(figsize=(12, 6))

# Plotting Training and Validation Accuracy
plt.subplot(1, 2, 1)
for i in range(n_splits):
    plt.plot(epochs, train_acc_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_acc_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Plotting Training and Validation Loss
plt.subplot(1, 2, 2)
for i in range(n_splits):
    plt.plot(epochs, train_loss_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_loss_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()

```



```
[10]: def build_model():
    base_model = VGG19(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

    # Freeze first 10 convolutional layers
    for layer in base_model.layers[:17]:
        layer.trainable = False

    x = base_model.output

    # Flatten to prepare for the fully connected layers
    x = Flatten()(x)

    # Fully connected + ReLU
    x = Dense(512, activation='relu')(x)
    x = Dense(1024, activation='relu')(x)

    # Adding output layer
    predictions = Dense(len(sheep_breed_list), activation='softmax')(x)

    # Create the final model
    model_VGG19 = Model(inputs=base_model.input, outputs=predictions)

    # Custom learning rates
    lr_mult = {}

    # Set learning rate to 0 for the first 10 layers
    for layer in model_VGG19.layers[:17]:
```

```

        lr_mult[layer.name + '/kernel:0'] = 0.0
        lr_mult[layer.name + '/bias:0'] = 0.0

    # Set learning rate to 0.0001 for the next Conv2D layers and the following
    ↪ layers
    for layer in model_VGG19.layers[17:23]:
        if isinstance(layer, Conv2D):
            lr_mult[layer.name + '/kernel:0'] = 0.001
            lr_mult[layer.name + '/bias:0'] = 0.001

    # Set learning rate to 10 for the last two dense layers
    for layer in model_VGG19.layers[23:]:
        lr_mult[layer.name + '/kernel:0'] = 0.001
        lr_mult[layer.name + '/bias:0'] = 0.001

    # Use SGD optimizer with initial learning rate of 1e-4
    optimizer = SGD(learning_rate=1e-4, momentum=0.9)

    # Compile the model
    model_VGG19.compile(optimizer=optimizer, loss='categorical_crossentropy',
    ↪ metrics=['accuracy'], loss_weights=[1.] + [10.]*2)

    return model_VGG19

accuracy_values = []
precision_values = []
recall_values = []
f1_values = []

train_acc_list = []
val_acc_list = []
train_loss_list = []
val_loss_list = []

# KFold Cross-validation
n_splits = 5
kf = KFold(n_splits=n_splits, shuffle=True, random_state=777)
fold_var = 1

for train_index, test_index in kf.split(img_data):
    X_train, X_test = img_data[train_index], img_data[test_index]
    y_train, y_test = labels_categorical[train_index],
    ↪ labels_categorical[test_index]

    model_VGG19 = build_model()

```

```

    #early_stop = EarlyStopping(monitor='val_loss', patience=3, verbose=1,
    ↪mode='min')
    save_path='/content/drive/MyDrive/Deep_learning_projects/SheepFace_Data/
    ↪SheepFace_three.keras'
    checkpoint =
    ↪ModelCheckpoint(save_path,monitor='val_loss',save_best_only=True,verbose=0)
    history = model_VGG19.fit(X_train, y_train, batch_size=10, epochs=10,
    ↪verbose=1, validation_split=0.2, callbacks=[checkpoint])

    print("\n")
    print(f'Evaluating the Test metrics')
    model_VGG19=load_model(save_path)
    # Evaluate the model
    scores = model_VGG19.evaluate(X_test, y_test, verbose=1)
    print(f'Test loss for fold {fold_var}: {scores[0]}')
    print(f'Test accuracy for fold {fold_var}: {scores[1]}')

    accuracy_values.append(scores[1]) # Appending accuracy to the list

    # Get predictions from the model
    y_true = np.argmax(y_test, axis=1) # Convert one-hot encoded labels back
    ↪to categorical labels
    y_pred = np.argmax(model_VGG19.predict(X_test), axis=1) # Get predictions
    ↪from the model

    # Calculate and append evaluation metrics for this fold
    precision = precision_score(y_true, y_pred, average='weighted',
    ↪zero_division=1)
    recall = recall_score(y_true, y_pred, average='weighted')
    f1 = f1_score(y_true, y_pred, average='weighted')

    precision_values.append(precision)
    recall_values.append(recall)
    f1_values.append(f1)

    # Print evaluation metrics for this fold
    print(f'Precision for fold {fold_var}: {precision}')
    print(f'Recall for fold {fold_var}: {recall}')
    print(f'F1 Score for fold {fold_var}: {f1}')
    print("\n")

    train_acc_list.append(history.history['accuracy'])
    val_acc_list.append(history.history['val_accuracy'])
    train_loss_list.append(history.history['loss'])
    val_loss_list.append(history.history['val_loss'])

```

```

# Increment the fold number
fold_var += 1

# Calculate standard deviation of accuracy across folds
accuracy_std_dev = np.std(accuracy_values)
print(f'Standard deviation of accuracy across folds: {accuracy_std_dev}')

# Calculate average of accuracy, precision, recall, and F1 across folds
avg_accuracy = np.mean(accuracy_values)
avg_precision = np.mean(precision_values)
avg_recall = np.mean(recall_values)
avg_f1 = np.mean(f1_values)

print(f'Average Accuracy across folds: {avg_accuracy}')
print(f'Average Precision across folds: {avg_precision}')
print(f'Average Recall across folds: {avg_recall}')
print(f'Average F1 Score across folds: {avg_f1}')

```

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5
80134624/80134624 [=====] - 3s 0us/step
Epoch 1/10
108/108 [=====] - 33s 295ms/step - loss: 0.9192 -
accuracy: 0.7293 - val_loss: 0.2874 - val_accuracy: 0.8736
Epoch 2/10
108/108 [=====] - 11s 103ms/step - loss: 0.0734 -
accuracy: 0.9786 - val_loss: 0.2228 - val_accuracy: 0.9071
Epoch 3/10
108/108 [=====] - 11s 103ms/step - loss: 0.0050 -
accuracy: 1.0000 - val_loss: 0.1990 - val_accuracy: 0.9294
Epoch 4/10
108/108 [=====] - 11s 104ms/step - loss: 0.0017 -
accuracy: 1.0000 - val_loss: 0.1823 - val_accuracy: 0.9368
Epoch 5/10
108/108 [=====] - 11s 105ms/step - loss: 0.0011 -
accuracy: 1.0000 - val_loss: 0.1803 - val_accuracy: 0.9368
Epoch 6/10
108/108 [=====] - 11s 106ms/step - loss: 8.3498e-04 -
accuracy: 1.0000 - val_loss: 0.1789 - val_accuracy: 0.9368
Epoch 7/10
108/108 [=====] - 11s 106ms/step - loss: 6.7482e-04 -
accuracy: 1.0000 - val_loss: 0.1785 - val_accuracy: 0.9368
Epoch 8/10
108/108 [=====] - 11s 106ms/step - loss: 5.6869e-04 -
accuracy: 1.0000 - val_loss: 0.1783 - val_accuracy: 0.9405
Epoch 9/10
108/108 [=====] - 11s 106ms/step - loss: 4.9131e-04 -

```



accuracy: 1.0000 - val\_loss: 0.1779 - val\_accuracy: 0.9405  
Epoch 10/10  
108/108 [=====] - 11s 106ms/step - loss: 4.3272e-04 -  
accuracy: 1.0000 - val\_loss: 0.1777 - val\_accuracy: 0.9405

Evaluating the Test metrics

11/11 [=====] - 2s 141ms/step - loss: 0.2327 -  
accuracy: 0.9345  
Test loss for fold 1: 0.23270520567893982  
Test accuracy for fold 1: 0.9345238208770752  
11/11 [=====] - 2s 145ms/step  
Precision for fold 1: 0.9347579126015797  
Recall for fold 1: 0.9345238095238095  
F1 Score for fold 1: 0.9338567012238453

Epoch 1/10  
108/108 [=====] - 13s 112ms/step - loss: 0.9492 -  
accuracy: 0.7153 - val\_loss: 0.2284 - val\_accuracy: 0.9145  
Epoch 2/10  
108/108 [=====] - 11s 106ms/step - loss: 0.0796 -  
accuracy: 0.9786 - val\_loss: 0.1523 - val\_accuracy: 0.9480  
Epoch 3/10  
108/108 [=====] - 11s 107ms/step - loss: 0.0082 -  
accuracy: 1.0000 - val\_loss: 0.1110 - val\_accuracy: 0.9740  
Epoch 4/10  
108/108 [=====] - 11s 107ms/step - loss: 0.0025 -  
accuracy: 1.0000 - val\_loss: 0.1068 - val\_accuracy: 0.9665  
Epoch 5/10  
108/108 [=====] - 11s 107ms/step - loss: 0.0016 -  
accuracy: 1.0000 - val\_loss: 0.1053 - val\_accuracy: 0.9665  
Epoch 6/10  
108/108 [=====] - 12s 107ms/step - loss: 0.0012 -  
accuracy: 1.0000 - val\_loss: 0.1051 - val\_accuracy: 0.9628  
Epoch 7/10  
108/108 [=====] - 12s 107ms/step - loss: 9.9825e-04 -  
accuracy: 1.0000 - val\_loss: 0.1048 - val\_accuracy: 0.9628  
Epoch 8/10  
108/108 [=====] - 12s 107ms/step - loss: 8.4112e-04 -  
accuracy: 1.0000 - val\_loss: 0.1040 - val\_accuracy: 0.9628  
Epoch 9/10  
108/108 [=====] - 10s 97ms/step - loss: 7.2585e-04 -  
accuracy: 1.0000 - val\_loss: 0.1042 - val\_accuracy: 0.9628  
Epoch 10/10  
108/108 [=====] - 11s 98ms/step - loss: 6.4048e-04 -  
accuracy: 1.0000 - val\_loss: 0.1042 - val\_accuracy: 0.9628

Evaluating the Test metrics

11/11 [=====] - 2s 147ms/step - loss: 0.1571 -

accuracy: 0.9494

Test loss for fold 2: 0.15712346136569977

Test accuracy for fold 2: 0.949404776096344

11/11 [=====] - 2s 153ms/step

Precision for fold 2: 0.9496049867884554

Recall for fold 2: 0.9494047619047619

F1 Score for fold 2: 0.948893698388513

Epoch 1/10

108/108 [=====] - 13s 113ms/step - loss: 1.1769 -

accuracy: 0.6353 - val\_loss: 0.3152 - val\_accuracy: 0.8625

Epoch 2/10

108/108 [=====] - 12s 107ms/step - loss: 0.1306 -

accuracy: 0.9581 - val\_loss: 0.1551 - val\_accuracy: 0.9480

Epoch 3/10

108/108 [=====] - 10s 97ms/step - loss: 0.0212 -

accuracy: 0.9953 - val\_loss: 0.1704 - val\_accuracy: 0.9405

Epoch 4/10

108/108 [=====] - 11s 106ms/step - loss: 0.0045 -

accuracy: 1.0000 - val\_loss: 0.1249 - val\_accuracy: 0.9480

Epoch 5/10

108/108 [=====] - 12s 107ms/step - loss: 0.0017 -

accuracy: 1.0000 - val\_loss: 0.1208 - val\_accuracy: 0.9517

Epoch 6/10

108/108 [=====] - 12s 107ms/step - loss: 0.0013 -

accuracy: 1.0000 - val\_loss: 0.1182 - val\_accuracy: 0.9554

Epoch 7/10

108/108 [=====] - 12s 108ms/step - loss: 9.8905e-04 -

accuracy: 1.0000 - val\_loss: 0.1161 - val\_accuracy: 0.9628

Epoch 8/10

108/108 [=====] - 12s 108ms/step - loss: 8.2004e-04 -

accuracy: 1.0000 - val\_loss: 0.1150 - val\_accuracy: 0.9591

Epoch 9/10

108/108 [=====] - 12s 109ms/step - loss: 7.0010e-04 -

accuracy: 1.0000 - val\_loss: 0.1137 - val\_accuracy: 0.9665

Epoch 10/10

108/108 [=====] - 12s 109ms/step - loss: 6.1486e-04 -

accuracy: 1.0000 - val\_loss: 0.1121 - val\_accuracy: 0.9628

Evaluating the Test metrics

11/11 [=====] - 2s 146ms/step - loss: 0.2644 -

accuracy: 0.9315

Test loss for fold 3: 0.26438698172569275

Test accuracy for fold 3: 0.9315476417541504  
11/11 [=====] - 2s 152ms/step  
Precision for fold 3: 0.9352569286588563  
Recall for fold 3: 0.9315476190476191  
F1 Score for fold 3: 0.9315647185936404

Epoch 1/10  
108/108 [=====] - 13s 112ms/step - loss: 1.1272 -  
accuracy: 0.6847 - val\_loss: 0.4102 - val\_accuracy: 0.8253  
Epoch 2/10  
108/108 [=====] - 12s 107ms/step - loss: 0.1148 -  
accuracy: 0.9647 - val\_loss: 0.2277 - val\_accuracy: 0.9294  
Epoch 3/10  
108/108 [=====] - 11s 107ms/step - loss: 0.0110 -  
accuracy: 1.0000 - val\_loss: 0.1858 - val\_accuracy: 0.9368  
Epoch 4/10  
108/108 [=====] - 12s 107ms/step - loss: 0.0025 -  
accuracy: 1.0000 - val\_loss: 0.1770 - val\_accuracy: 0.9331  
Epoch 5/10  
108/108 [=====] - 12s 107ms/step - loss: 0.0016 -  
accuracy: 1.0000 - val\_loss: 0.1725 - val\_accuracy: 0.9331  
Epoch 6/10  
108/108 [=====] - 12s 108ms/step - loss: 0.0012 -  
accuracy: 1.0000 - val\_loss: 0.1713 - val\_accuracy: 0.9294  
Epoch 7/10  
108/108 [=====] - 12s 109ms/step - loss: 9.7238e-04 -  
accuracy: 1.0000 - val\_loss: 0.1687 - val\_accuracy: 0.9294  
Epoch 8/10  
108/108 [=====] - 11s 100ms/step - loss: 8.1678e-04 -  
accuracy: 1.0000 - val\_loss: 0.1688 - val\_accuracy: 0.9219  
Epoch 9/10  
108/108 [=====] - 12s 109ms/step - loss: 7.0584e-04 -  
accuracy: 1.0000 - val\_loss: 0.1670 - val\_accuracy: 0.9219  
Epoch 10/10  
108/108 [=====] - 12s 109ms/step - loss: 6.2270e-04 -  
accuracy: 1.0000 - val\_loss: 0.1668 - val\_accuracy: 0.9219

Evaluating the Test metrics  
11/11 [=====] - 2s 145ms/step - loss: 0.1404 -  
accuracy: 0.9405  
Test loss for fold 4: 0.14040616154670715  
Test accuracy for fold 4: 0.9404761791229248  
11/11 [=====] - 2s 148ms/step  
Precision for fold 4: 0.9408698287927612  
Recall for fold 4: 0.9404761904761905  
F1 Score for fold 4: 0.940162089054884

Epoch 1/10  
108/108 [=====] - 13s 112ms/step - loss: 0.9651 -  
accuracy: 0.7265 - val\_loss: 0.3392 - val\_accuracy: 0.8736  
Epoch 2/10  
108/108 [=====] - 12s 107ms/step - loss: 0.0807 -  
accuracy: 0.9758 - val\_loss: 0.1698 - val\_accuracy: 0.9554  
Epoch 3/10  
108/108 [=====] - 12s 107ms/step - loss: 0.0191 -  
accuracy: 0.9963 - val\_loss: 0.1469 - val\_accuracy: 0.9517  
Epoch 4/10  
108/108 [=====] - 12s 107ms/step - loss: 0.0021 -  
accuracy: 1.0000 - val\_loss: 0.1454 - val\_accuracy: 0.9517  
Epoch 5/10  
108/108 [=====] - 12s 108ms/step - loss: 0.0012 -  
accuracy: 1.0000 - val\_loss: 0.1404 - val\_accuracy: 0.9554  
Epoch 6/10  
108/108 [=====] - 12s 108ms/step - loss: 9.3455e-04 -  
accuracy: 1.0000 - val\_loss: 0.1388 - val\_accuracy: 0.9554  
Epoch 7/10  
108/108 [=====] - 12s 109ms/step - loss: 7.5677e-04 -  
accuracy: 1.0000 - val\_loss: 0.1373 - val\_accuracy: 0.9554  
Epoch 8/10  
108/108 [=====] - 12s 109ms/step - loss: 6.4048e-04 -  
accuracy: 1.0000 - val\_loss: 0.1369 - val\_accuracy: 0.9591  
Epoch 9/10  
108/108 [=====] - 12s 109ms/step - loss: 5.5299e-04 -  
accuracy: 1.0000 - val\_loss: 0.1359 - val\_accuracy: 0.9591  
Epoch 10/10  
108/108 [=====] - 12s 109ms/step - loss: 4.8711e-04 -  
accuracy: 1.0000 - val\_loss: 0.1351 - val\_accuracy: 0.9591

Evaluating the Test metrics

11/11 [=====] - 2s 147ms/step - loss: 0.2572 -  
accuracy: 0.9375  
Test loss for fold 5: 0.2572050988674164  
Test accuracy for fold 5: 0.9375  
11/11 [=====] - 2s 152ms/step  
Precision for fold 5: 0.93798679769608  
Recall for fold 5: 0.9375  
F1 Score for fold 5: 0.9370556523004459

Standard deviation of accuracy across folds: 0.006128348738136682  
Average Accuracy across folds: 0.9386904835700989  
Average Precision across folds: 0.9396952909075467

Average Recall across folds: 0.9386904761904763  
Average F1 Score across folds: 0.9383065719122656

```
[11]: import matplotlib.pyplot as plt

# Determine the maximum length among all lists
max_length = max(len(train_acc) for train_acc in train_acc_list)

# Pad the shorter lists with zeros to match the maximum length
for i in range(n_splits):
    train_acc_list[i] += [0] * (max_length - len(train_acc_list[i]))
    val_acc_list[i] += [0] * (max_length - len(val_acc_list[i]))
    train_loss_list[i] += [0] * (max_length - len(train_loss_list[i]))
    val_loss_list[i] += [0] * (max_length - len(val_loss_list[i]))


# Plotting
epochs = range(1, max_length + 1)

plt.figure(figsize=(12, 6))

# Plotting Training and Validation Accuracy
plt.subplot(1, 2, 1)
for i in range(n_splits):
    plt.plot(epochs, train_acc_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_acc_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Plotting Training and Validation Loss
plt.subplot(1, 2, 2)
for i in range(n_splits):
    plt.plot(epochs, train_loss_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_loss_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()
```



```
[6]: from keras.applications.vgg16 import VGG16, preprocess_input
from sklearn.svm import SVC
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import StratifiedKFold
import numpy as np

encoder = LabelEncoder()
labels_encoded = encoder.fit_transform(labels)

# Load the pre-trained VGG16 model without the top (classification) layers
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze the layers of the pre-trained model
for layer in base_model.layers:
    layer.trainable = False

# Define SVM pipeline
svm_pipeline = Pipeline([('scaler', StandardScaler()), ('svm', SVC(kernel='linear', probability=True))])

# Initialize lists to store evaluation metrics
accuracy_values = []
precision_values = []
```

```

recall_values = []
f1_values = []

# Perform k-fold cross-validation
n_splits = 5
skf = StratifiedKFold(n_splits=n_splits, shuffle=True, random_state=777)
fold_var = 1

for train_index, test_index in skf.split(img_data, labels_encoded): # Ensure
    ↪ labels_encoded is defined correctly as the non-categorical labels
    X_train, X_test = img_data[train_index], img_data[test_index]
    y_train, y_test = labels_encoded[train_index], labels_encoded[test_index]

    # Extract features using the VGG16 model and flatten them
    X_train_features = base_model.predict(preprocess_input(X_train))
    X_train_features = X_train_features.reshape(X_train_features.shape[0], -1)

    X_test_features = base_model.predict(preprocess_input(X_test))
    X_test_features = X_test_features.reshape(X_test_features.shape[0], -1)

    # Fit and transform the training data with StandardScaler and train the SVM
    svm_pipeline.fit(X_train_features, y_train)

    # Predict using the trained SVM pipeline
    predictions = svm_pipeline.predict(X_test_features)
    print(f'Fold {fold_var}:')
    print(f'Accuracy: {accuracy_score(y_test, predictions)}')
    print(f'Precision: {precision_score(y_test, predictions,
    ↪ average="weighted")}')
    print(f'Recall: {recall_score(y_test, predictions, average="weighted")}')
    print(f'F1 Score: {f1_score(y_test, predictions, average="weighted")}')

    # Calculate evaluation metrics
    accuracy_values.append(accuracy_score(y_test, predictions))
    precision_values.append(precision_score(y_test, predictions,
    ↪ average='weighted'))
    recall_values.append(recall_score(y_test, predictions, average='weighted'))
    f1_values.append(f1_score(y_test, predictions, average='weighted'))

    # Increment the fold number
    fold_var += 1

# Calculate average metrics across all folds
avg_accuracy = np.mean(accuracy_values)
avg_precision = np.mean(precision_values)
avg_recall = np.mean(recall_values)
avg_f1 = np.mean(f1_values)

```

```

print(f'Average test Accuracy across folds: {avg_accuracy}')
print(f'Average Precision across folds: {avg_precision}')
print(f'Average Recall across folds: {avg_recall}')
print(f'Average F1 Score across folds: {avg_f1}')

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)  
 58889256/58889256 [=====] - 3s 0us/step

42/42 [=====] - 12s 108ms/step

11/11 [=====] - 5s 474ms/step

Fold 1:

Accuracy: 0.9553571428571429

Precision: 0.9558990724255652

Recall: 0.9553571428571429

F1 Score: 0.9548037603627753

42/42 [=====] - 5s 118ms/step

11/11 [=====] - 1s 116ms/step

Fold 2:

Accuracy: 0.9702380952380952

Precision: 0.9702761627906978

Recall: 0.9702380952380952

F1 Score: 0.9700502152080345

42/42 [=====] - 5s 121ms/step

11/11 [=====] - 1s 120ms/step

Fold 3:

Accuracy: 0.9732142857142857

Precision: 0.9734078992890944

Recall: 0.9732142857142857

F1 Score: 0.9731926934849087

42/42 [=====] - 5s 125ms/step

11/11 [=====] - 1s 125ms/step

Fold 4:

Accuracy: 0.9642857142857143

Precision: 0.965298202614379

Recall: 0.9642857142857143

F1 Score: 0.9642785203194726

42/42 [=====] - 5s 130ms/step

11/11 [=====] - 1s 128ms/step

Fold 5:

Accuracy: 0.9791666666666666

Precision: 0.9793767507002801

Recall: 0.9791666666666666

F1 Score: 0.9792188619227544

Average test Accuracy across folds: 0.9684523809523811

Average Precision across folds: 0.9688516175640034

Average Recall across folds: 0.9684523809523811



Average F1 Score across folds: 0.968308810259589

```
[12]: !jupyter nbconvert --to pdf '/content/drive/MyDrive/Colab_Notebooks/  
      ↪VGG19_Sheep_Classification.ipynb'
```

```
[NbConvertApp] Converting notebook  
/content/drive/MyDrive/Colab_Notebooks/VGG19_Sheep_Classification.ipynb to pdf  
[NbConvertApp] Support files will be in VGG19_Sheep_Classification_files/  
[NbConvertApp] Making directory ./VGG19_Sheep_Classification_files  
[NbConvertApp] Making directory ./VGG19_Sheep_Classification_files  
[NbConvertApp] Making directory ./VGG19_Sheep_Classification_files  
[NbConvertApp] Writing 153313 bytes to notebook.tex  
[NbConvertApp] Building PDF  
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']  
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']  
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no  
citations  
[NbConvertApp] PDF successfully created  
[NbConvertApp] Writing 476732 bytes to  
/content/drive/MyDrive/Colab_Notebooks/VGG19_Sheep_Classification.pdf
```