

# Sheep\_Full\_Images\_Classification

April 28, 2024

```
[6]: import numpy as np
import pandas as pd
import os
from PIL import Image
from keras.applications.vgg16 import VGG16, preprocess_input
from tensorflow.keras.applications import VGG19
from keras.preprocessing import image
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Model, load_model
from keras.layers import Dense, GlobalAveragePooling2D, Dropout,
    ↳BatchNormalization, Flatten, Conv2D, Activation, MaxPooling2D
from keras.callbacks import EarlyStopping, Callback, ModelCheckpoint
from sklearn.model_selection import train_test_split, KFold
from tensorflow.keras.utils import to_categorical
import tensorflow as tf
from tensorflow.keras.optimizers import SGD
from sklearn.utils import shuffle
from sklearn.model_selection import StratifiedKFold
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↳f1_score
```

```
[9]: # Load images and labels
path = '/content/drive/MyDrive/Deep_learning_projects/SheepFullImages'
sheep_breed_list = os.listdir(path)
print(sheep_breed_list)
```

```
['Suffolk', 'Poll Dorset', 'White Suffolk', 'Marino']
```

```
[10]: def load_and_preprocess_images(path, target_size=(224, 224)):
    img_data_list = []
    labels = []

    for idx, sheep_breed in enumerate(sheep_breed_list):
        sheep_breed_files = os.listdir(os.path.join(path, sheep_breed))
        print(sheep_breed_files)
        for img_file in sheep_breed_files:
            sheep_image_path = os.path.join(path, sheep_breed, img_file)
```

```

        try:
            img = image.load_img(sheep_image_path, target_size=target_size)
            img_array = image.img_to_array(img)
            img_array = preprocess_input(img_array)
            img_data_list.append(img_array)
            labels.append(idx)
        except Exception as e:
            print(f"Error loading image {img_file}: {e}")

    # Convert lists to numpy arrays and preprocess
    img_data = np.array(img_data_list)
    labels = np.array(labels)

    # Shuffle data
    img_data, labels = shuffle(img_data, labels, random_state=777)

    return img_data, labels, sheep_breed_list

img_data_path = '/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/
↳sheep_img_data.npy'
labels_path = '/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/
↳sheep_labels.npy'

if os.path.exists(img_data_path) and os.path.exists(labels_path):
    img_data = np.load(img_data_path)
    labels = np.load(labels_path)
else:
    img_data, labels, sheep_breed_list = load_and_preprocess_images(path)
    # Save the processed data
    np.save(img_data_path, img_data)
    np.save(labels_path, labels)

print('Data Shape:', img_data.shape)
print('Labels Shape:', labels.shape)

```

Data Shape: (1619, 224, 224, 3)  
Labels Shape: (1619,)

```

[11]: # One-hot encode labels
labels_categorical = to_categorical(labels, num_classes=len(sheep_breed_list))

print(labels_categorical.shape)

```

(1619, 4)

```

[12]: def build_model():

```

```

base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze first 10 convolutional layers
for layer in base_model.layers[:15]:
    layer.trainable = False

x = base_model.output

# Flatten to prepare for the fully connected layers
x = Flatten()(x)

# Fully connected + ReLU
x = Dense(64, activation='relu')(x)
x = Dense(64, activation='relu')(x)

# Adding output layer
predictions = Dense(len(sheep_breed_list), activation='softmax')(x)

# Create the final model
model = Model(inputs=base_model.input, outputs=predictions)

# Custom learning rates
lr_mult = {}

# Set learning rate to 0 for the first 10 layers
for layer in model.layers[:15]:
    lr_mult[layer.name + '/kernel:0'] = 0.0
    lr_mult[layer.name + '/bias:0'] = 0.0

# Set learning rate to 0.0001 for the next Conv2D layers and the following layers
for layer in model.layers[15:-4]:
    if isinstance(layer, Conv2D):
        lr_mult[layer.name + '/kernel:0'] = 0.0001
        lr_mult[layer.name + '/bias:0'] = 0.0001

# Set learning rate to 10 for the last two dense layers
for layer in model.layers[-4:]:
    lr_mult[layer.name + '/kernel:0'] = 10
    lr_mult[layer.name + '/bias:0'] = 10

# Use SGD optimizer with initial learning rate of 1e-4
optimizer = SGD(learning_rate=1e-4, momentum=0.9)

# Compile the model

```

```

        model.compile(optimizer=optimizer, loss='categorical_crossentropy',
↪metrics=['accuracy'], loss_weights=[1.] + [10.]*2)

        return model

accuracy_values = []
precision_values = []
recall_values = []
f1_values = []

train_acc_list = []
val_acc_list = []
train_loss_list = []
val_loss_list = []

# KFold Cross-validation
n_splits = 5
kf = KFold(n_splits=n_splits, shuffle=True, random_state=777)
fold_var = 1

for train_index, test_index in kf.split(img_data):
    X_train, X_test = img_data[train_index], img_data[test_index]
    y_train, y_test = labels_categorical[train_index],
↪labels_categorical[test_index]

    model = build_model()

    #early_stop = EarlyStopping(monitor='val_loss', patience=3, verbose=1,
↪mode='min')
    save_path='/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/
↪sheepFace_one.h5'
    checkpoint =
↪ModelCheckpoint(save_path,monitor='val_loss',save_best_only=True,verbose=1)
    history = model.fit(X_train, y_train, batch_size=10, epochs=10, verbose=1,
↪validation_split=0.2, callbacks=[checkpoint])

    print("\n")
    print(f'Evaluating the Test metrics')
    model=load_model(save_path)
    # Evaluate the model
    scores = model.evaluate(X_test, y_test, verbose=1)
    print(f'Test loss for fold {fold_var}: {scores[0]}')
    print(f'Test accuracy for fold {fold_var}: {scores[1]}')

    accuracy_values.append(scores[1]) # Appending accuracy to the list

```

```

    # Get predictions from the model
    y_true = np.argmax(y_test, axis=1) # Convert one-hot encoded labels back
    to categorical labels
    y_pred = np.argmax(model.predict(X_test), axis=1) # Get predictions from
    the model

    # Calculate and append evaluation metrics for this fold
    precision = precision_score(y_true, y_pred, average='weighted',
    zero_division=1)
    recall = recall_score(y_true, y_pred, average='weighted')
    f1 = f1_score(y_true, y_pred, average='weighted')

    precision_values.append(precision)
    recall_values.append(recall)
    f1_values.append(f1)

    # Print evaluation metrics for this fold
    print(f'Precision for fold {fold_var}: {precision}')
    print(f'Recall for fold {fold_var}: {recall}')
    print(f'F1 Score for fold {fold_var}: {f1}')
    print("\n")

    train_acc_list.append(history.history['accuracy'])
    val_acc_list.append(history.history['val_accuracy'])
    train_loss_list.append(history.history['loss'])
    val_loss_list.append(history.history['val_loss'])

    # Increment the fold number
    fold_var += 1

# Calculate standard deviation of accuracy across folds
accuracy_std_dev = np.std(accuracy_values)
print(f'Standard deviation of accuracy across folds: {accuracy_std_dev}')

# Calculate average of accuracy, precision, recall, and F1 across folds
avg_accuracy = np.mean(accuracy_values)
avg_precision = np.mean(precision_values)
avg_recall = np.mean(recall_values)
avg_f1 = np.mean(f1_values)

print(f'Average Accuracy across folds: {avg_accuracy}')
print(f'Average Precision across folds: {avg_precision}')
print(f'Average Recall across folds: {avg_recall}')
print(f'Average F1 Score across folds: {avg_f1}')

```

Epoch 1/10

104/104 [=====] - ETA: 0s - loss: 1.5677 - accuracy:

```

0.5502
Epoch 1: val_loss improved from inf to 0.93505, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/sheepFace_one.h5
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:
UserWarning: You are saving your model as an HDF5 file via `model.save()`. This
file format is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')`.
    saving_api.save_model(

104/104 [=====] - 29s 218ms/step - loss: 1.5677 -
accuracy: 0.5502 - val_loss: 0.9350 - val_accuracy: 0.5792
Epoch 2/10
104/104 [=====] - ETA: 0s - loss: 0.5933 - accuracy:
0.7683
Epoch 2: val_loss improved from 0.93505 to 0.49222, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/sheepFace_one.h5
104/104 [=====] - 9s 84ms/step - loss: 0.5933 -
accuracy: 0.7683 - val_loss: 0.4922 - val_accuracy: 0.7876
Epoch 3/10
104/104 [=====] - ETA: 0s - loss: 0.3133 - accuracy:
0.8851
Epoch 3: val_loss improved from 0.49222 to 0.41848, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/sheepFace_one.h5
104/104 [=====] - 10s 95ms/step - loss: 0.3133 -
accuracy: 0.8851 - val_loss: 0.4185 - val_accuracy: 0.8610
Epoch 4/10
104/104 [=====] - ETA: 0s - loss: 0.1248 - accuracy:
0.9595
Epoch 4: val_loss improved from 0.41848 to 0.29627, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/sheepFace_one.h5
104/104 [=====] - 9s 89ms/step - loss: 0.1248 -
accuracy: 0.9595 - val_loss: 0.2963 - val_accuracy: 0.9112
Epoch 5/10
104/104 [=====] - ETA: 0s - loss: 0.0681 - accuracy:
0.9797
Epoch 5: val_loss did not improve from 0.29627
104/104 [=====] - 9s 83ms/step - loss: 0.0681 -
accuracy: 0.9797 - val_loss: 0.3818 - val_accuracy: 0.8842
Epoch 6/10
103/104 [=====>.] - ETA: 0s - loss: 0.0374 - accuracy:
0.9893
Epoch 6: val_loss improved from 0.29627 to 0.25930, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/sheepFace_one.h5
104/104 [=====] - 10s 93ms/step - loss: 0.0372 -
accuracy: 0.9894 - val_loss: 0.2593 - val_accuracy: 0.9189
Epoch 7/10
103/104 [=====>.] - ETA: 0s - loss: 0.0136 - accuracy:
0.9981

```

Epoch 7: val\_loss improved from 0.25930 to 0.23679, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 82ms/step - loss: 0.0136 -  
accuracy: 0.9981 - val\_loss: 0.2368 - val\_accuracy: 0.9151  
Epoch 8/10  
103/104 [=====>.] - ETA: 0s - loss: 0.0056 - accuracy:  
1.0000  
Epoch 8: val\_loss improved from 0.23679 to 0.21585, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 82ms/step - loss: 0.0056 -  
accuracy: 1.0000 - val\_loss: 0.2158 - val\_accuracy: 0.9305  
Epoch 9/10  
103/104 [=====>.] - ETA: 0s - loss: 0.0032 - accuracy:  
1.0000  
Epoch 9: val\_loss did not improve from 0.21585  
104/104 [=====] - 8s 79ms/step - loss: 0.0032 -  
accuracy: 1.0000 - val\_loss: 0.2173 - val\_accuracy: 0.9228  
Epoch 10/10  
103/104 [=====>.] - ETA: 0s - loss: 0.0024 - accuracy:  
1.0000  
Epoch 10: val\_loss did not improve from 0.21585  
104/104 [=====] - 8s 79ms/step - loss: 0.0024 -  
accuracy: 1.0000 - val\_loss: 0.2197 - val\_accuracy: 0.9228

#### Evaluating the Test metrics

11/11 [=====] - 10s 295ms/step - loss: 0.3883 -  
accuracy: 0.8827  
Test loss for fold 1: 0.3883349597454071  
Test accuracy for fold 1: 0.8827160596847534  
11/11 [=====] - 2s 137ms/step  
Precision for fold 1: 0.885663930768558  
Recall for fold 1: 0.8827160493827161  
F1 Score for fold 1: 0.883510431785993

Epoch 1/10  
104/104 [=====] - ETA: 0s - loss: 1.6973 - accuracy:  
0.5145  
Epoch 1: val\_loss improved from inf to 0.86636, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:  
UserWarning: You are saving your model as an HDF5 file via `model.save()`. This  
file format is considered legacy. We recommend using instead the native Keras  
format, e.g. `model.save('my\_model.keras')`.  
saving\_api.save\_model(  
104/104 [=====] - 10s 90ms/step - loss: 1.6973 -

accuracy: 0.5145 - val\_loss: 0.8664 - val\_accuracy: 0.6448  
Epoch 2/10  
104/104 [=====] - ETA: 0s - loss: 0.6605 - accuracy: 0.7249  
Epoch 2: val\_loss improved from 0.86636 to 0.51221, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 86ms/step - loss: 0.6605 - accuracy: 0.7249 - val\_loss: 0.5122 - val\_accuracy: 0.8108  
Epoch 3/10  
104/104 [=====] - ETA: 0s - loss: 0.3179 - accuracy: 0.8861  
Epoch 3: val\_loss improved from 0.51221 to 0.38019, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 10s 97ms/step - loss: 0.3179 - accuracy: 0.8861 - val\_loss: 0.3802 - val\_accuracy: 0.8687  
Epoch 4/10  
104/104 [=====] - ETA: 0s - loss: 0.1594 - accuracy: 0.9488  
Epoch 4: val\_loss improved from 0.38019 to 0.32078, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 85ms/step - loss: 0.1594 - accuracy: 0.9488 - val\_loss: 0.3208 - val\_accuracy: 0.8996  
Epoch 5/10  
104/104 [=====] - ETA: 0s - loss: 0.0806 - accuracy: 0.9788  
Epoch 5: val\_loss did not improve from 0.32078  
104/104 [=====] - 8s 81ms/step - loss: 0.0806 - accuracy: 0.9788 - val\_loss: 0.3937 - val\_accuracy: 0.8764  
Epoch 6/10  
103/104 [=====>.] - ETA: 0s - loss: 0.0436 - accuracy: 0.9845  
Epoch 6: val\_loss improved from 0.32078 to 0.29129, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 91ms/step - loss: 0.0435 - accuracy: 0.9846 - val\_loss: 0.2913 - val\_accuracy: 0.9035  
Epoch 7/10  
103/104 [=====>.] - ETA: 0s - loss: 0.0171 - accuracy: 0.9990  
Epoch 7: val\_loss improved from 0.29129 to 0.28810, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 83ms/step - loss: 0.0173 - accuracy: 0.9990 - val\_loss: 0.2881 - val\_accuracy: 0.9073  
Epoch 8/10  
103/104 [=====>.] - ETA: 0s - loss: 0.0071 - accuracy: 0.9990  
Epoch 8: val\_loss improved from 0.28810 to 0.24419, saving model to /content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 83ms/step - loss: 0.0071 -



```

accuracy: 0.9990 - val_loss: 0.2442 - val_accuracy: 0.9035
Epoch 9/10
104/104 [=====] - ETA: 0s - loss: 0.0035 - accuracy:
1.0000
Epoch 9: val_loss did not improve from 0.24419
104/104 [=====] - 8s 80ms/step - loss: 0.0035 -
accuracy: 1.0000 - val_loss: 0.2603 - val_accuracy: 0.9073
Epoch 10/10
104/104 [=====] - ETA: 0s - loss: 0.0027 - accuracy:
1.0000
Epoch 10: val_loss did not improve from 0.24419
104/104 [=====] - 8s 81ms/step - loss: 0.0027 -
accuracy: 1.0000 - val_loss: 0.2639 - val_accuracy: 0.9073

```

#### Evaluating the Test metrics

```

11/11 [=====] - 2s 125ms/step - loss: 0.2831 -
accuracy: 0.9228
Test loss for fold 2: 0.28305599093437195
Test accuracy for fold 2: 0.9228395223617554
11/11 [=====] - 2s 135ms/step
Precision for fold 2: 0.9236707633133129
Recall for fold 2: 0.9228395061728395
F1 Score for fold 2: 0.9228414232212712

```

```

Epoch 1/10
103/104 [=====>.] - ETA: 0s - loss: 2.0800 - accuracy:
0.4194
Epoch 1: val_loss improved from inf to 1.07769, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/sheepFace_one.h5
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:
UserWarning: You are saving your model as an HDF5 file via `model.save()`. This
file format is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')`.
    saving_api.save_model(
104/104 [=====] - 10s 88ms/step - loss: 2.0730 -
accuracy: 0.4208 - val_loss: 1.0777 - val_accuracy: 0.5637
Epoch 2/10
104/104 [=====] - ETA: 0s - loss: 0.8500 - accuracy:
0.6525
Epoch 2: val_loss improved from 1.07769 to 0.79991, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/sheepFace_one.h5
104/104 [=====] - 9s 84ms/step - loss: 0.8500 -
accuracy: 0.6525 - val_loss: 0.7999 - val_accuracy: 0.6873
Epoch 3/10
104/104 [=====] - ETA: 0s - loss: 0.6054 - accuracy:

```

0.7654  
Epoch 3: val\_loss improved from 0.79991 to 0.52569, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 85ms/step - loss: 0.6054 -  
accuracy: 0.7654 - val\_loss: 0.5257 - val\_accuracy: 0.8069  
Epoch 4/10  
104/104 [=====] - ETA: 0s - loss: 0.4074 - accuracy:  
0.8446  
Epoch 4: val\_loss improved from 0.52569 to 0.45125, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 84ms/step - loss: 0.4074 -  
accuracy: 0.8446 - val\_loss: 0.4512 - val\_accuracy: 0.8378  
Epoch 5/10  
104/104 [=====] - ETA: 0s - loss: 0.2413 - accuracy:  
0.9122  
Epoch 5: val\_loss improved from 0.45125 to 0.33636, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 84ms/step - loss: 0.2413 -  
accuracy: 0.9122 - val\_loss: 0.3364 - val\_accuracy: 0.8842  
Epoch 6/10  
104/104 [=====] - ETA: 0s - loss: 0.1420 - accuracy:  
0.9537  
Epoch 6: val\_loss improved from 0.33636 to 0.27814, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 83ms/step - loss: 0.1420 -  
accuracy: 0.9537 - val\_loss: 0.2781 - val\_accuracy: 0.9151  
Epoch 7/10  
103/104 [=====>.] - ETA: 0s - loss: 0.0698 - accuracy:  
0.9786  
Epoch 7: val\_loss improved from 0.27814 to 0.23493, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
104/104 [=====] - 9s 83ms/step - loss: 0.0695 -  
accuracy: 0.9788 - val\_loss: 0.2349 - val\_accuracy: 0.9112  
Epoch 8/10  
104/104 [=====] - ETA: 0s - loss: 0.0360 - accuracy:  
0.9942  
Epoch 8: val\_loss did not improve from 0.23493  
104/104 [=====] - 8s 80ms/step - loss: 0.0360 -  
accuracy: 0.9942 - val\_loss: 0.2950 - val\_accuracy: 0.9151  
Epoch 9/10  
103/104 [=====>.] - ETA: 0s - loss: 0.0229 - accuracy:  
0.9961  
Epoch 9: val\_loss did not improve from 0.23493  
104/104 [=====] - 8s 80ms/step - loss: 0.0231 -  
accuracy: 0.9961 - val\_loss: 0.2548 - val\_accuracy: 0.9344  
Epoch 10/10  
103/104 [=====>.] - ETA: 0s - loss: 0.0102 - accuracy:  
1.0000

Epoch 10: val\_loss did not improve from 0.23493  
104/104 [=====] - 8s 80ms/step - loss: 0.0101 -  
accuracy: 1.0000 - val\_loss: 0.2448 - val\_accuracy: 0.9421

Evaluating the Test metrics

11/11 [=====] - 2s 127ms/step - loss: 0.2914 -  
accuracy: 0.8951  
Test loss for fold 3: 0.29137712717056274  
Test accuracy for fold 3: 0.895061731338501  
11/11 [=====] - 1s 133ms/step  
Precision for fold 3: 0.8962645225192823  
Recall for fold 3: 0.8950617283950617  
F1 Score for fold 3: 0.8950502549577888

Epoch 1/10

104/104 [=====] - ETA: 0s - loss: 1.9542 - accuracy:  
0.4614  
Epoch 1: val\_loss improved from inf to 0.92482, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5  
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:  
UserWarning: You are saving your model as an HDF5 file via `model.save()`. This  
file format is considered legacy. We recommend using instead the native Keras  
format, e.g. `model.save('my\_model.keras')`.  
saving\_api.save\_model(

104/104 [=====] - 10s 88ms/step - loss: 1.9542 -  
accuracy: 0.4614 - val\_loss: 0.9248 - val\_accuracy: 0.6139

Epoch 2/10

104/104 [=====] - ETA: 0s - loss: 0.7298 - accuracy:  
0.7037

Epoch 2: val\_loss improved from 0.92482 to 0.74546, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5

104/104 [=====] - 9s 84ms/step - loss: 0.7298 -  
accuracy: 0.7037 - val\_loss: 0.7455 - val\_accuracy: 0.7027

Epoch 3/10

104/104 [=====] - ETA: 0s - loss: 0.3575 - accuracy:  
0.8639

Epoch 3: val\_loss improved from 0.74546 to 0.44737, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5

104/104 [=====] - 9s 85ms/step - loss: 0.3575 -  
accuracy: 0.8639 - val\_loss: 0.4474 - val\_accuracy: 0.8224

Epoch 4/10

104/104 [=====] - ETA: 0s - loss: 0.1788 - accuracy:  
0.9431

Epoch 4: val\_loss did not improve from 0.44737

104/104 [=====] - 9s 82ms/step - loss: 0.1788 -

```

accuracy: 0.9431 - val_loss: 0.5606 - val_accuracy: 0.7838
Epoch 5/10
103/104 [=====>.] - ETA: 0s - loss: 0.1307 - accuracy:
0.9544
Epoch 5: val_loss did not improve from 0.44737
104/104 [=====] - 8s 81ms/step - loss: 0.1301 -
accuracy: 0.9546 - val_loss: 0.4535 - val_accuracy: 0.8378
Epoch 6/10
104/104 [=====] - ETA: 0s - loss: 0.0710 - accuracy:
0.9720
Epoch 6: val_loss improved from 0.44737 to 0.22405, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/sheepFace_one.h5
104/104 [=====] - 9s 91ms/step - loss: 0.0710 -
accuracy: 0.9720 - val_loss: 0.2241 - val_accuracy: 0.9305
Epoch 7/10
103/104 [=====>.] - ETA: 0s - loss: 0.0295 - accuracy:
0.9922
Epoch 7: val_loss did not improve from 0.22405
104/104 [=====] - 8s 80ms/step - loss: 0.0294 -
accuracy: 0.9923 - val_loss: 0.2351 - val_accuracy: 0.9266
Epoch 8/10
103/104 [=====>.] - ETA: 0s - loss: 0.0101 - accuracy:
0.9990
Epoch 8: val_loss improved from 0.22405 to 0.21597, saving model to
/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/sheepFace_one.h5
104/104 [=====] - 9s 88ms/step - loss: 0.0102 -
accuracy: 0.9990 - val_loss: 0.2160 - val_accuracy: 0.9344
Epoch 9/10
103/104 [=====>.] - ETA: 0s - loss: 0.0116 - accuracy:
0.9981
Epoch 9: val_loss did not improve from 0.21597
104/104 [=====] - 8s 80ms/step - loss: 0.0115 -
accuracy: 0.9981 - val_loss: 0.2353 - val_accuracy: 0.9112
Epoch 10/10
104/104 [=====] - ETA: 0s - loss: 0.0046 - accuracy:
1.0000
Epoch 10: val_loss did not improve from 0.21597
104/104 [=====] - 8s 80ms/step - loss: 0.0046 -
accuracy: 1.0000 - val_loss: 0.2213 - val_accuracy: 0.9266

```

#### Evaluating the Test metrics

```

11/11 [=====] - 2s 127ms/step - loss: 0.2731 -
accuracy: 0.9167
Test loss for fold 4: 0.2730903923511505
Test accuracy for fold 4: 0.9166666865348816
11/11 [=====] - 1s 134ms/step
Precision for fold 4: 0.9176818682115054

```

Recall for fold 4: 0.9166666666666666  
F1 Score for fold 4: 0.9168440981220994

Epoch 1/10

103/104 [=====>.] - ETA: 0s - loss: 1.6284 - accuracy: 0.5544

Epoch 1: val\_loss improved from inf to 0.84469, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5

/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:

UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')`.

saving\_api.save\_model(

104/104 [=====] - 10s 89ms/step - loss: 1.6251 -  
accuracy: 0.5550 - val\_loss: 0.8447 - val\_accuracy: 0.6462

Epoch 2/10

104/104 [=====] - ETA: 0s - loss: 0.5358 - accuracy: 0.8021

Epoch 2: val\_loss improved from 0.84469 to 0.51912, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5

104/104 [=====] - 9s 84ms/step - loss: 0.5358 -  
accuracy: 0.8021 - val\_loss: 0.5191 - val\_accuracy: 0.8038

Epoch 3/10

104/104 [=====] - ETA: 0s - loss: 0.2590 - accuracy: 0.9102

Epoch 3: val\_loss improved from 0.51912 to 0.36596, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5

104/104 [=====] - 9s 86ms/step - loss: 0.2590 -  
accuracy: 0.9102 - val\_loss: 0.3660 - val\_accuracy: 0.8769

Epoch 4/10

103/104 [=====>.] - ETA: 0s - loss: 0.0956 - accuracy: 0.9757

Epoch 4: val\_loss did not improve from 0.36596

104/104 [=====] - 9s 83ms/step - loss: 0.0951 -  
accuracy: 0.9759 - val\_loss: 0.3802 - val\_accuracy: 0.8654

Epoch 5/10

104/104 [=====] - ETA: 0s - loss: 0.0504 - accuracy: 0.9865

Epoch 5: val\_loss improved from 0.36596 to 0.30576, saving model to  
/content/drive/MyDrive/Deep\_learning\_projects/SheepFull\_Data/sheepFace\_one.h5

104/104 [=====] - 9s 90ms/step - loss: 0.0504 -  
accuracy: 0.9865 - val\_loss: 0.3058 - val\_accuracy: 0.8923

Epoch 6/10

103/104 [=====>.] - ETA: 0s - loss: 0.0156 - accuracy: 0.9981

Epoch 6: val\_loss improved from 0.30576 to 0.24608, saving model to

```

/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/sheepFace_one.h5
104/104 [=====] - 9s 83ms/step - loss: 0.0155 -
accuracy: 0.9981 - val_loss: 0.2461 - val_accuracy: 0.9115
Epoch 7/10
103/104 [=====>.] - ETA: 0s - loss: 0.0061 - accuracy:
1.0000
Epoch 7: val_loss did not improve from 0.24608
104/104 [=====] - 8s 80ms/step - loss: 0.0061 -
accuracy: 1.0000 - val_loss: 0.2833 - val_accuracy: 0.9038
Epoch 8/10
103/104 [=====>.] - ETA: 0s - loss: 0.0030 - accuracy:
1.0000
Epoch 8: val_loss did not improve from 0.24608
104/104 [=====] - 8s 80ms/step - loss: 0.0030 -
accuracy: 1.0000 - val_loss: 0.2858 - val_accuracy: 0.9115
Epoch 9/10
103/104 [=====>.] - ETA: 0s - loss: 0.0023 - accuracy:
1.0000
Epoch 9: val_loss did not improve from 0.24608
104/104 [=====] - 8s 80ms/step - loss: 0.0023 -
accuracy: 1.0000 - val_loss: 0.2973 - val_accuracy: 0.9077
Epoch 10/10
103/104 [=====>.] - ETA: 0s - loss: 0.0018 - accuracy:
1.0000
Epoch 10: val_loss did not improve from 0.24608
104/104 [=====] - 8s 81ms/step - loss: 0.0018 -
accuracy: 1.0000 - val_loss: 0.3048 - val_accuracy: 0.9077

```

#### Evaluating the Test metrics

```

11/11 [=====] - 3s 273ms/step - loss: 0.1921 -
accuracy: 0.9195
Test loss for fold 5: 0.192056804895401
Test accuracy for fold 5: 0.9195046424865723
11/11 [=====] - 2s 135ms/step
Precision for fold 5: 0.9195274331859169
Recall for fold 5: 0.9195046439628483
F1 Score for fold 5: 0.9192853929696035

```

```

Standard deviation of accuracy across folds: 0.01569900346178792
Average Accuracy across folds: 0.9073577284812927
Average Precision across folds: 0.9085617035997151
Average Recall across folds: 0.9073577189160265
Average F1 Score across folds: 0.9075063202113511

```

```
[15]: import matplotlib.pyplot as plt

# Determine the maximum length among all lists
max_length = max(len(train_acc) for train_acc in train_acc_list)

# Pad the shorter lists with zeros to match the maximum length
for i in range(n_splits):
    train_acc_list[i] += [0] * (max_length - len(train_acc_list[i]))
    val_acc_list[i] += [0] * (max_length - len(val_acc_list[i]))
    train_loss_list[i] += [0] * (max_length - len(train_loss_list[i]))
    val_loss_list[i] += [0] * (max_length - len(val_loss_list[i]))

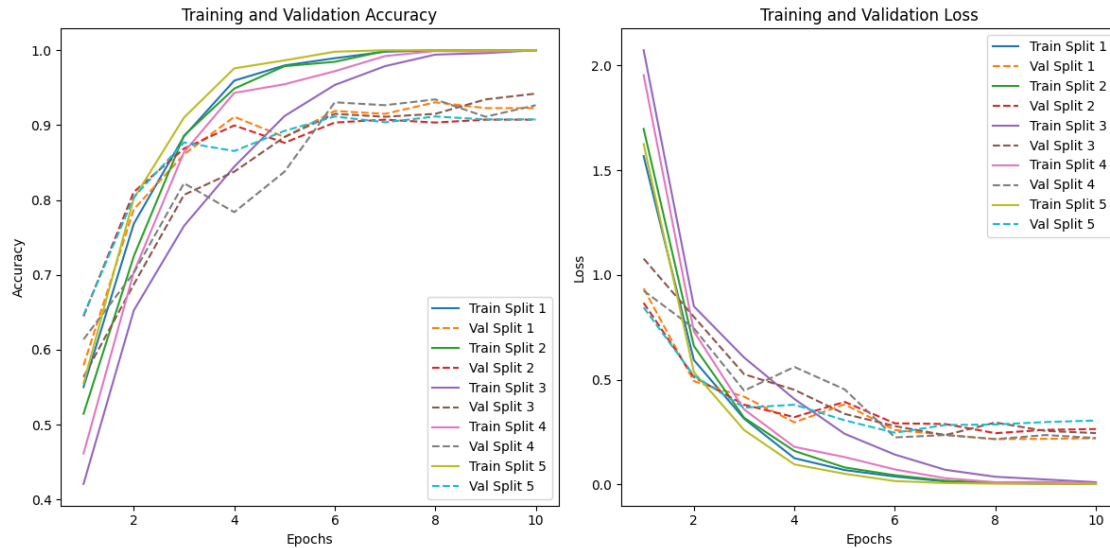
# Plotting
epochs = range(1, max_length + 1)

plt.figure(figsize=(12, 6))

# Plotting Training and Validation Accuracy
plt.subplot(1, 2, 1)
for i in range(n_splits):
    plt.plot(epochs, train_acc_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_acc_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Plotting Training and Validation Loss
plt.subplot(1, 2, 2)
for i in range(n_splits):
    plt.plot(epochs, train_loss_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_loss_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()
```



```
[11]: def build_model():
    base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

    # Freeze first 10 convolutional layers
    for layer in base_model.layers[:15]:
        layer.trainable = False

    x = base_model.output

    # Adding Conv2D + ReLU Layers
    x = Conv2D(512, (3, 3), padding='same')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = Conv2D(512, (3, 3), padding='same', activation='relu')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = Conv2D(512, (3, 3), padding='same', activation='relu')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)

    # Adding Max Pooling
    x = MaxPooling2D((2, 2), strides=(2, 2))(x)

    # Add GlobalAveragePooling2D layer
    x = GlobalAveragePooling2D()(x)

    # Replacing the last three layers with new fully connected layers
```



```

x = Dense(256)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)

x = Dense(256)(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Dropout(0.5)(x)

predictions = Dense(len(sheep_breed_list), activation='softmax')(x) # Add
↳ output layer

model_normalizeddd = Model(inputs=base_model.input, outputs=predictions)

# Custom learning rates
lr_mult = {}

# Set learning rate to 0 for the first 10 layers
for layer in model_normalizeddd.layers[:15]:
    lr_mult[layer.name + '/kernel:0'] = 0.0
    lr_mult[layer.name + '/bias:0'] = 0.0

# Set learning rate to 0.0001 for the next Conv2D layers and the following
↳ layers
for layer in model_normalizeddd.layers[15:30]:
    if isinstance(layer, Conv2D):
        lr_mult[layer.name + '/kernel:0'] = 0.0001
        lr_mult[layer.name + '/bias:0'] = 0.0001

# Set learning rate to 10 for the last two dense layers
for layer in model_normalizeddd.layers[30:]:
    lr_mult[layer.name + '/kernel:0'] = 1
    lr_mult[layer.name + '/bias:0'] = 1

# Use SGD optimizer with initial learning rate of 1e-4
optimizer = SGD(learning_rate=1e-4, momentum=0.9)

# Compile the model
model_normalizeddd.compile(optimizer=optimizer,
↳ loss='categorical_crossentropy', metrics=['accuracy'], loss_weights=[1.] +
↳ [10.]*2)

return model_normalizeddd

accuracy_values = []
precision_values = []

```

```

recall_values = []
f1_values = []

train_acc_list = []
val_acc_list = []
train_loss_list = []
val_loss_list = []

# KFold Cross-validation
n_splits = 5
kf = KFold(n_splits=n_splits, shuffle=True, random_state=777)
fold_var = 1

for train_index, test_index in kf.split(img_data):
    X_train, X_test = img_data[train_index], img_data[test_index]
    y_train, y_test = labels_categorical[train_index],
    ↪ labels_categorical[test_index]

    model_normalizedd = build_model()

    # Define the ImageDataGenerator for data augmentation
    datagen = ImageDataGenerator(
        width_shift_range=0.1, # translate horizontally by 10% of total width
        height_shift_range=0.1, # translate vertically by 10% of total height
        fill_mode='nearest' # strategy for filling in newly created pixels
    )

    # Generate augmented images batches during training
    train_datagen = datagen.flow(X_train, y_train, batch_size=10)

    #early_stop = EarlyStopping(monitor='val_loss', patience=3, verbose=1, ↪
    ↪ mode='min')
    save_path = '/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/
    ↪ sheepFace_two.keras'
    checkpoint =
    ↪ ModelCheckpoint(save_path, monitor='val_loss', save_best_only=True, verbose=0)
    history = model_normalizedd.fit(train_datagen, epochs=20,
    ↪ validation_data=(X_test, y_test), verbose=1, callbacks=[checkpoint])

    print("\n")
    print(f'Evaluating the Test metrics')
    model_normalizedd = load_model(save_path)
    # Evaluate the model
    scores = model_normalizedd.evaluate(X_test, y_test, verbose=1)
    print(f'Test loss for fold {fold_var}: {scores[0]}')
    print(f'Test accuracy for fold {fold_var}: {scores[1]}')

```

```

accuracy_values.append(scores[1]) # Appending accuracy to the list

# Get predictions from the model
y_true = np.argmax(y_test, axis=1) # Convert one-hot encoded labels back
↳to categorical labels
y_pred = np.argmax(model_normalizedd.predict(X_test), axis=1) # Get
↳predictions from the model

# Calculate and append evaluation metrics for this fold
precision = precision_score(y_true, y_pred, average='weighted',
↳zero_division=1)
recall = recall_score(y_true, y_pred, average='weighted')
f1 = f1_score(y_true, y_pred, average='weighted')

precision_values.append(precision)
recall_values.append(recall)
f1_values.append(f1)

# Print evaluation metrics for this fold
print(f'Precision for fold {fold_var}: {precision}')
print(f'Recall for fold {fold_var}: {recall}')
print(f'F1 Score for fold {fold_var}: {f1}')

train_acc_list.append(history.history['accuracy'])
val_acc_list.append(history.history['val_accuracy'])
train_loss_list.append(history.history['loss'])
val_loss_list.append(history.history['val_loss'])
print("\n")

# Increment the fold number
fold_var += 1

# Calculate standard deviation of accuracy across folds
accuracy_std_dev = np.std(accuracy_values)
print(f'Standard deviation of accuracy across folds: {accuracy_std_dev}')

# Calculate average of accuracy, precision, recall, and F1 across folds
avg_accuracy = np.mean(accuracy_values)
avg_precision = np.mean(precision_values)
avg_recall = np.mean(recall_values)
avg_f1 = np.mean(f1_values)

print(f'Average Accuracy across folds: {avg_accuracy}')
print(f'Average Precision across folds: {avg_precision}')
print(f'Average Recall across folds: {avg_recall}')
print(f'Average F1 Score across folds: {avg_f1}')

```

Epoch 1/20  
130/130 [=====] - 24s 146ms/step - loss: 1.7989 - accuracy: 0.3220 - val\_loss: 1.2810 - val\_accuracy: 0.3827

Epoch 2/20  
130/130 [=====] - 16s 120ms/step - loss: 1.2778 - accuracy: 0.4587 - val\_loss: 1.0687 - val\_accuracy: 0.5093

Epoch 3/20  
130/130 [=====] - 15s 118ms/step - loss: 1.0691 - accuracy: 0.5761 - val\_loss: 0.7457 - val\_accuracy: 0.7377

Epoch 4/20  
130/130 [=====] - 16s 119ms/step - loss: 0.8291 - accuracy: 0.6587 - val\_loss: 0.5276 - val\_accuracy: 0.8210

Epoch 5/20  
130/130 [=====] - 15s 114ms/step - loss: 0.6811 - accuracy: 0.7421 - val\_loss: 0.5913 - val\_accuracy: 0.7840

Epoch 6/20  
130/130 [=====] - 15s 118ms/step - loss: 0.5906 - accuracy: 0.7838 - val\_loss: 0.3870 - val\_accuracy: 0.8642

Epoch 7/20  
130/130 [=====] - 16s 119ms/step - loss: 0.5173 - accuracy: 0.8139 - val\_loss: 0.3757 - val\_accuracy: 0.8673

Epoch 8/20  
130/130 [=====] - 15s 115ms/step - loss: 0.4585 - accuracy: 0.8363 - val\_loss: 0.5545 - val\_accuracy: 0.8056

Epoch 9/20  
130/130 [=====] - 15s 118ms/step - loss: 0.4209 - accuracy: 0.8471 - val\_loss: 0.2577 - val\_accuracy: 0.9259

Epoch 10/20  
130/130 [=====] - 15s 114ms/step - loss: 0.4046 - accuracy: 0.8602 - val\_loss: 0.4228 - val\_accuracy: 0.8673

Epoch 11/20  
130/130 [=====] - 14s 111ms/step - loss: 0.3067 - accuracy: 0.8927 - val\_loss: 0.3194 - val\_accuracy: 0.9012

Epoch 12/20  
130/130 [=====] - 15s 118ms/step - loss: 0.2787 - accuracy: 0.9097 - val\_loss: 0.1761 - val\_accuracy: 0.9537

Epoch 13/20  
130/130 [=====] - 15s 113ms/step - loss: 0.2792 - accuracy: 0.9120 - val\_loss: 0.1818 - val\_accuracy: 0.9321

Epoch 14/20  
130/130 [=====] - 14s 110ms/step - loss: 0.2458 - accuracy: 0.9158 - val\_loss: 0.3486 - val\_accuracy: 0.8827

Epoch 15/20  
130/130 [=====] - 15s 111ms/step - loss: 0.2472 - accuracy: 0.9266 - val\_loss: 0.2110 - val\_accuracy: 0.9383

Epoch 16/20  
130/130 [=====] - 15s 114ms/step - loss: 0.2637 - accuracy: 0.9189 - val\_loss: 0.2293 - val\_accuracy: 0.9228

Epoch 17/20  
130/130 [=====] - 15s 112ms/step - loss: 0.2059 -  
accuracy: 0.9375 - val\_loss: 0.1787 - val\_accuracy: 0.9506  
Epoch 18/20  
130/130 [=====] - 14s 110ms/step - loss: 0.1843 -  
accuracy: 0.9483 - val\_loss: 0.1783 - val\_accuracy: 0.9568  
Epoch 19/20  
130/130 [=====] - 15s 111ms/step - loss: 0.1959 -  
accuracy: 0.9405 - val\_loss: 0.2128 - val\_accuracy: 0.9321  
Epoch 20/20  
130/130 [=====] - 15s 117ms/step - loss: 0.1712 -  
accuracy: 0.9467 - val\_loss: 0.1443 - val\_accuracy: 0.9475

#### Evaluating the Test metrics

11/11 [=====] - 2s 132ms/step - loss: 0.1443 -  
accuracy: 0.9475  
Test loss for fold 1: 0.1442643404006958  
Test accuracy for fold 1: 0.9475308656692505  
11/11 [=====] - 2s 141ms/step  
Precision for fold 1: 0.9485916873803986  
Recall for fold 1: 0.9475308641975309  
F1 Score for fold 1: 0.9472657986768309

Epoch 1/20  
130/130 [=====] - 18s 121ms/step - loss: 1.6908 -  
accuracy: 0.3452 - val\_loss: 1.1815 - val\_accuracy: 0.5185  
Epoch 2/20  
130/130 [=====] - 15s 116ms/step - loss: 1.2419 -  
accuracy: 0.4996 - val\_loss: 0.8404 - val\_accuracy: 0.6451  
Epoch 3/20  
130/130 [=====] - 15s 118ms/step - loss: 0.9764 -  
accuracy: 0.6224 - val\_loss: 0.6378 - val\_accuracy: 0.7870  
Epoch 4/20  
130/130 [=====] - 15s 117ms/step - loss: 0.7801 -  
accuracy: 0.6996 - val\_loss: 0.4306 - val\_accuracy: 0.8457  
Epoch 5/20  
130/130 [=====] - 15s 117ms/step - loss: 0.6128 -  
accuracy: 0.7568 - val\_loss: 0.3875 - val\_accuracy: 0.8580  
Epoch 6/20  
130/130 [=====] - 16s 119ms/step - loss: 0.5599 -  
accuracy: 0.7915 - val\_loss: 0.2982 - val\_accuracy: 0.8981  
Epoch 7/20  
130/130 [=====] - 16s 120ms/step - loss: 0.4627 -  
accuracy: 0.8309 - val\_loss: 0.2729 - val\_accuracy: 0.9167  
Epoch 8/20  
130/130 [=====] - 16s 120ms/step - loss: 0.4165 -

```

accuracy: 0.8602 - val_loss: 0.2565 - val_accuracy: 0.9198
Epoch 9/20
130/130 [=====] - 15s 113ms/step - loss: 0.4071 -
accuracy: 0.8649 - val_loss: 0.4176 - val_accuracy: 0.8457
Epoch 10/20
130/130 [=====] - 15s 118ms/step - loss: 0.3641 -
accuracy: 0.8764 - val_loss: 0.2443 - val_accuracy: 0.9105
Epoch 11/20
130/130 [=====] - 16s 120ms/step - loss: 0.3183 -
accuracy: 0.8919 - val_loss: 0.1918 - val_accuracy: 0.9444
Epoch 12/20
130/130 [=====] - 15s 117ms/step - loss: 0.3177 -
accuracy: 0.8896 - val_loss: 0.1871 - val_accuracy: 0.9352
Epoch 13/20
130/130 [=====] - 15s 116ms/step - loss: 0.2619 -
accuracy: 0.9243 - val_loss: 0.1730 - val_accuracy: 0.9568
Epoch 14/20
130/130 [=====] - 15s 112ms/step - loss: 0.2515 -
accuracy: 0.9166 - val_loss: 0.2513 - val_accuracy: 0.9105
Epoch 15/20
130/130 [=====] - 15s 113ms/step - loss: 0.2557 -
accuracy: 0.9166 - val_loss: 0.1984 - val_accuracy: 0.9475
Epoch 16/20
130/130 [=====] - 14s 111ms/step - loss: 0.2162 -
accuracy: 0.9297 - val_loss: 0.2119 - val_accuracy: 0.9352
Epoch 17/20
130/130 [=====] - 15s 115ms/step - loss: 0.1816 -
accuracy: 0.9459 - val_loss: 0.1426 - val_accuracy: 0.9506
Epoch 18/20
130/130 [=====] - 15s 118ms/step - loss: 0.2266 -
accuracy: 0.9282 - val_loss: 0.1416 - val_accuracy: 0.9537
Epoch 19/20
130/130 [=====] - 15s 118ms/step - loss: 0.1728 -
accuracy: 0.9483 - val_loss: 0.1248 - val_accuracy: 0.9537
Epoch 20/20
130/130 [=====] - 15s 113ms/step - loss: 0.2038 -
accuracy: 0.9413 - val_loss: 0.2455 - val_accuracy: 0.9290

```

#### Evaluating the Test metrics

```

11/11 [=====] - 2s 131ms/step - loss: 0.1248 -
accuracy: 0.9537
Test loss for fold 2: 0.12477420270442963
Test accuracy for fold 2: 0.9537037014961243
11/11 [=====] - 2s 141ms/step
Precision for fold 2: 0.9537671286921809
Recall for fold 2: 0.9537037037037037
F1 Score for fold 2: 0.9536984533043342

```

Epoch 1/20  
130/130 [=====] - 18s 123ms/step - loss: 1.7451 - accuracy: 0.3328 - val\_loss: 1.4398 - val\_accuracy: 0.3765

Epoch 2/20  
130/130 [=====] - 15s 118ms/step - loss: 1.3072 - accuracy: 0.4819 - val\_loss: 0.8803 - val\_accuracy: 0.5957

Epoch 3/20  
130/130 [=====] - 15s 118ms/step - loss: 1.0268 - accuracy: 0.5892 - val\_loss: 0.6872 - val\_accuracy: 0.7315

Epoch 4/20  
130/130 [=====] - 16s 120ms/step - loss: 0.8356 - accuracy: 0.6811 - val\_loss: 0.4653 - val\_accuracy: 0.8673

Epoch 5/20  
130/130 [=====] - 16s 120ms/step - loss: 0.6616 - accuracy: 0.7313 - val\_loss: 0.3635 - val\_accuracy: 0.8765

Epoch 6/20  
130/130 [=====] - 15s 118ms/step - loss: 0.5258 - accuracy: 0.8023 - val\_loss: 0.3396 - val\_accuracy: 0.8642

Epoch 7/20  
130/130 [=====] - 15s 112ms/step - loss: 0.5094 - accuracy: 0.8170 - val\_loss: 0.6713 - val\_accuracy: 0.7623

Epoch 8/20  
130/130 [=====] - 15s 112ms/step - loss: 0.4224 - accuracy: 0.8525 - val\_loss: 0.3464 - val\_accuracy: 0.8611

Epoch 9/20  
130/130 [=====] - 15s 117ms/step - loss: 0.3778 - accuracy: 0.8595 - val\_loss: 0.2020 - val\_accuracy: 0.9228

Epoch 10/20  
130/130 [=====] - 15s 113ms/step - loss: 0.3311 - accuracy: 0.8849 - val\_loss: 0.2830 - val\_accuracy: 0.9012

Epoch 11/20  
130/130 [=====] - 15s 111ms/step - loss: 0.2922 - accuracy: 0.9012 - val\_loss: 0.2803 - val\_accuracy: 0.8951

Epoch 12/20  
130/130 [=====] - 15s 111ms/step - loss: 0.3015 - accuracy: 0.9027 - val\_loss: 0.2312 - val\_accuracy: 0.9074

Epoch 13/20  
130/130 [=====] - 15s 114ms/step - loss: 0.2444 - accuracy: 0.9220 - val\_loss: 0.2469 - val\_accuracy: 0.9105

Epoch 14/20  
130/130 [=====] - 15s 118ms/step - loss: 0.2490 - accuracy: 0.9166 - val\_loss: 0.1782 - val\_accuracy: 0.9383

Epoch 15/20  
130/130 [=====] - 15s 114ms/step - loss: 0.2286 - accuracy: 0.9259 - val\_loss: 0.2394 - val\_accuracy: 0.9074

Epoch 16/20

130/130 [=====] - 14s 110ms/step - loss: 0.2167 -  
accuracy: 0.9313 - val\_loss: 0.3007 - val\_accuracy: 0.8951  
Epoch 17/20  
130/130 [=====] - 15s 112ms/step - loss: 0.1966 -  
accuracy: 0.9467 - val\_loss: 0.2711 - val\_accuracy: 0.9074  
Epoch 18/20  
130/130 [=====] - 15s 117ms/step - loss: 0.2054 -  
accuracy: 0.9367 - val\_loss: 0.1777 - val\_accuracy: 0.9475  
Epoch 19/20  
130/130 [=====] - 15s 114ms/step - loss: 0.2099 -  
accuracy: 0.9359 - val\_loss: 0.3097 - val\_accuracy: 0.8920  
Epoch 20/20  
130/130 [=====] - 14s 110ms/step - loss: 0.1629 -  
accuracy: 0.9552 - val\_loss: 0.1895 - val\_accuracy: 0.9444

#### Evaluating the Test metrics

11/11 [=====] - 2s 132ms/step - loss: 0.1777 -  
accuracy: 0.9475  
Test loss for fold 3: 0.1777118742465973  
Test accuracy for fold 3: 0.9475308656692505  
11/11 [=====] - 2s 141ms/step  
Precision for fold 3: 0.951475860330027  
Recall for fold 3: 0.9475308641975309  
F1 Score for fold 3: 0.947713761208653

Epoch 1/20  
130/130 [=====] - 19s 121ms/step - loss: 1.5454 -  
accuracy: 0.3753 - val\_loss: 1.7065 - val\_accuracy: 0.3519  
Epoch 2/20  
130/130 [=====] - 16s 119ms/step - loss: 1.2504 -  
accuracy: 0.4880 - val\_loss: 1.0285 - val\_accuracy: 0.5648  
Epoch 3/20  
130/130 [=====] - 15s 117ms/step - loss: 0.9940 -  
accuracy: 0.5931 - val\_loss: 0.7108 - val\_accuracy: 0.6944  
Epoch 4/20  
130/130 [=====] - 15s 118ms/step - loss: 0.8390 -  
accuracy: 0.6625 - val\_loss: 0.5412 - val\_accuracy: 0.8056  
Epoch 5/20  
130/130 [=====] - 15s 118ms/step - loss: 0.6787 -  
accuracy: 0.7282 - val\_loss: 0.4743 - val\_accuracy: 0.8241  
Epoch 6/20  
130/130 [=====] - 15s 117ms/step - loss: 0.5811 -  
accuracy: 0.7830 - val\_loss: 0.2881 - val\_accuracy: 0.9074  
Epoch 7/20  
130/130 [=====] - 15s 111ms/step - loss: 0.5089 -  
accuracy: 0.8255 - val\_loss: 0.3594 - val\_accuracy: 0.8796



Epoch 8/20  
130/130 [=====] - 15s 119ms/step - loss: 0.4549 - accuracy: 0.8378 - val\_loss: 0.2355 - val\_accuracy: 0.9167  
Epoch 9/20  
130/130 [=====] - 16s 119ms/step - loss: 0.4253 - accuracy: 0.8463 - val\_loss: 0.2027 - val\_accuracy: 0.9475  
Epoch 10/20  
130/130 [=====] - 15s 117ms/step - loss: 0.3932 - accuracy: 0.8579 - val\_loss: 0.1719 - val\_accuracy: 0.9537  
Epoch 11/20  
130/130 [=====] - 14s 110ms/step - loss: 0.3530 - accuracy: 0.8764 - val\_loss: 0.2438 - val\_accuracy: 0.9167  
Epoch 12/20  
130/130 [=====] - 15s 112ms/step - loss: 0.3104 - accuracy: 0.9035 - val\_loss: 0.1917 - val\_accuracy: 0.9352  
Epoch 13/20  
130/130 [=====] - 15s 115ms/step - loss: 0.2967 - accuracy: 0.9027 - val\_loss: 0.1484 - val\_accuracy: 0.9568  
Epoch 14/20  
130/130 [=====] - 15s 112ms/step - loss: 0.2738 - accuracy: 0.9135 - val\_loss: 0.1562 - val\_accuracy: 0.9475  
Epoch 15/20  
130/130 [=====] - 15s 117ms/step - loss: 0.2885 - accuracy: 0.9035 - val\_loss: 0.1459 - val\_accuracy: 0.9537  
Epoch 16/20  
130/130 [=====] - 15s 113ms/step - loss: 0.2519 - accuracy: 0.9181 - val\_loss: 0.1509 - val\_accuracy: 0.9506  
Epoch 17/20  
130/130 [=====] - 15s 117ms/step - loss: 0.2636 - accuracy: 0.9143 - val\_loss: 0.1104 - val\_accuracy: 0.9691  
Epoch 18/20  
130/130 [=====] - 15s 117ms/step - loss: 0.1862 - accuracy: 0.9398 - val\_loss: 0.1086 - val\_accuracy: 0.9660  
Epoch 19/20  
130/130 [=====] - 15s 118ms/step - loss: 0.1978 - accuracy: 0.9429 - val\_loss: 0.0909 - val\_accuracy: 0.9691  
Epoch 20/20  
130/130 [=====] - 15s 114ms/step - loss: 0.1954 - accuracy: 0.9382 - val\_loss: 0.1214 - val\_accuracy: 0.9506

Evaluating the Test metrics

11/11 [=====] - 2s 130ms/step - loss: 0.0909 - accuracy: 0.9691  
Test loss for fold 4: 0.09088943898677826  
Test accuracy for fold 4: 0.9691358208656311  
11/11 [=====] - 2s 139ms/step  
Precision for fold 4: 0.9700717250918054

Recall for fold 4: 0.9691358024691358  
F1 Score for fold 4: 0.9692712878650681

Epoch 1/20  
130/130 [=====] - 19s 126ms/step - loss: 1.6624 -  
accuracy: 0.3488 - val\_loss: 1.2553 - val\_accuracy: 0.4056  
Epoch 2/20  
130/130 [=====] - 16s 121ms/step - loss: 1.2464 -  
accuracy: 0.4985 - val\_loss: 0.8642 - val\_accuracy: 0.6625  
Epoch 3/20  
130/130 [=====] - 15s 118ms/step - loss: 0.9862 -  
accuracy: 0.6011 - val\_loss: 0.5330 - val\_accuracy: 0.8142  
Epoch 4/20  
130/130 [=====] - 15s 117ms/step - loss: 0.8055 -  
accuracy: 0.6798 - val\_loss: 0.3830 - val\_accuracy: 0.8576  
Epoch 5/20  
130/130 [=====] - 15s 118ms/step - loss: 0.6612 -  
accuracy: 0.7415 - val\_loss: 0.3250 - val\_accuracy: 0.8762  
Epoch 6/20  
130/130 [=====] - 16s 119ms/step - loss: 0.5472 -  
accuracy: 0.8009 - val\_loss: 0.2941 - val\_accuracy: 0.8885  
Epoch 7/20  
130/130 [=====] - 15s 118ms/step - loss: 0.4592 -  
accuracy: 0.8387 - val\_loss: 0.2613 - val\_accuracy: 0.8885  
Epoch 8/20  
130/130 [=====] - 15s 117ms/step - loss: 0.4070 -  
accuracy: 0.8596 - val\_loss: 0.1884 - val\_accuracy: 0.9350  
Epoch 9/20  
130/130 [=====] - 14s 110ms/step - loss: 0.3688 -  
accuracy: 0.8719 - val\_loss: 0.2418 - val\_accuracy: 0.9164  
Epoch 10/20  
130/130 [=====] - 15s 119ms/step - loss: 0.3401 -  
accuracy: 0.8935 - val\_loss: 0.1851 - val\_accuracy: 0.9443  
Epoch 11/20  
130/130 [=====] - 15s 119ms/step - loss: 0.2952 -  
accuracy: 0.8927 - val\_loss: 0.1689 - val\_accuracy: 0.9412  
Epoch 12/20  
130/130 [=====] - 15s 114ms/step - loss: 0.2971 -  
accuracy: 0.9020 - val\_loss: 0.1788 - val\_accuracy: 0.9350  
Epoch 13/20  
130/130 [=====] - 15s 111ms/step - loss: 0.2646 -  
accuracy: 0.9174 - val\_loss: 0.2108 - val\_accuracy: 0.9195  
Epoch 14/20  
130/130 [=====] - 15s 119ms/step - loss: 0.2432 -  
accuracy: 0.9267 - val\_loss: 0.1453 - val\_accuracy: 0.9505  
Epoch 15/20  
130/130 [=====] - 15s 113ms/step - loss: 0.2126 -

```

accuracy: 0.9313 - val_loss: 0.2594 - val_accuracy: 0.9102
Epoch 16/20
130/130 [=====] - 15s 112ms/step - loss: 0.2189 -
accuracy: 0.9306 - val_loss: 0.2209 - val_accuracy: 0.9288
Epoch 17/20
130/130 [=====] - 15s 112ms/step - loss: 0.1967 -
accuracy: 0.9398 - val_loss: 0.1937 - val_accuracy: 0.9288
Epoch 18/20
130/130 [=====] - 15s 113ms/step - loss: 0.1958 -
accuracy: 0.9406 - val_loss: 0.1550 - val_accuracy: 0.9659
Epoch 19/20
130/130 [=====] - 15s 116ms/step - loss: 0.1739 -
accuracy: 0.9560 - val_loss: 0.1395 - val_accuracy: 0.9505
Epoch 20/20
130/130 [=====] - 15s 113ms/step - loss: 0.1839 -
accuracy: 0.9421 - val_loss: 0.2541 - val_accuracy: 0.9071

```

Evaluating the Test metrics

```

11/11 [=====] - 2s 131ms/step - loss: 0.1395 -
accuracy: 0.9505
Test loss for fold 5: 0.1395476907491684
Test accuracy for fold 5: 0.9504643678665161
11/11 [=====] - 2s 141ms/step
Precision for fold 5: 0.9515184076329587
Recall for fold 5: 0.9504643962848297
F1 Score for fold 5: 0.9502868449010309

```

```

Standard deviation of accuracy across folds: 0.0080603563205253
Average Accuracy across folds: 0.9536731243133545
Average Precision across folds: 0.9550849618254741
Average Recall across folds: 0.9536731261705462
Average F1 Score across folds: 0.9536472291911835

```

```

[16]: import matplotlib.pyplot as plt

# Determine the maximum length among all lists
max_length = max(len(train_acc) for train_acc in train_acc_list)

# Pad the shorter lists with zeros to match the maximum length
for i in range(n_splits):
    train_acc_list[i] += [0] * (max_length - len(train_acc_list[i]))
    val_acc_list[i] += [0] * (max_length - len(val_acc_list[i]))
    train_loss_list[i] += [0] * (max_length - len(train_loss_list[i]))
    val_loss_list[i] += [0] * (max_length - len(val_loss_list[i]))

```

```

# Plotting
epochs = range(1, max_length + 1)

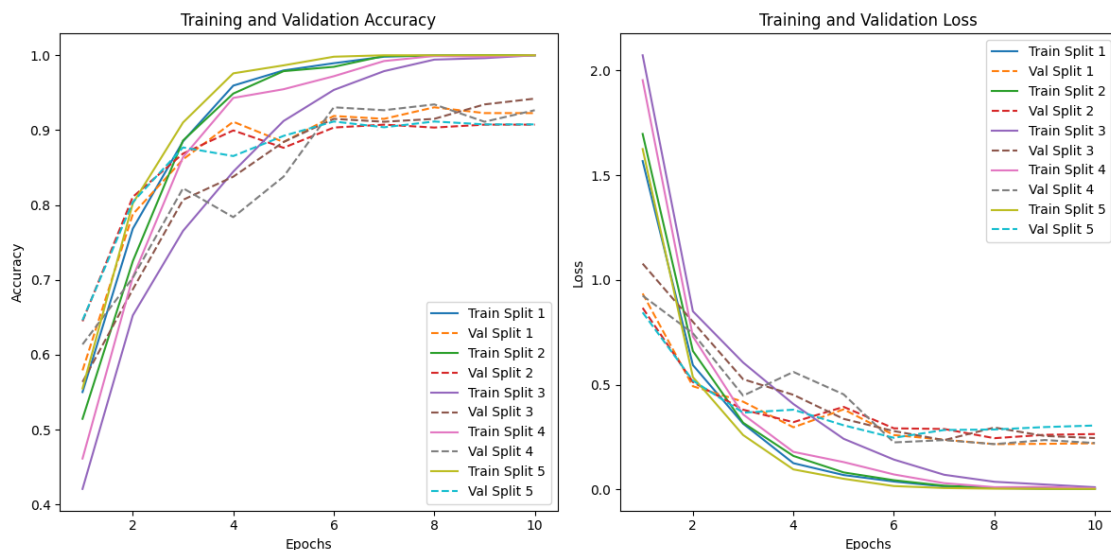
plt.figure(figsize=(12, 6))

# Plotting Training and Validation Accuracy
plt.subplot(1, 2, 1)
for i in range(n_splits):
    plt.plot(epochs, train_acc_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_acc_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Plotting Training and Validation Loss
plt.subplot(1, 2, 2)
for i in range(n_splits):
    plt.plot(epochs, train_loss_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_loss_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()

```



```
[8]: def build_model():
    base_model = VGG19(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

    # Freeze first 10 convolutional layers
    for layer in base_model.layers[:17]:
        layer.trainable = False

    x = base_model.output

    # Flatten to prepare for the fully connected layers
    x = Flatten()(x)

    # Fully connected + ReLU
    x = Dense(512, activation='relu')(x)
    x = Dense(1024, activation='relu')(x)

    # Adding output layer
    predictions = Dense(len(sheep_breed_list), activation='softmax')(x)

    # Create the final model
    model_VGG19 = Model(inputs=base_model.input, outputs=predictions)

    # Custom learning rates
    lr_mult = {}

    # Set learning rate to 0 for the first 10 layers
    for layer in model_VGG19.layers[:17]:
        lr_mult[layer.name + '/kernel:0'] = 0.0
        lr_mult[layer.name + '/bias:0'] = 0.0

    # Set learning rate to 0.0001 for the next Conv2D layers and the following layers
    for layer in model_VGG19.layers[17:23]:
        if isinstance(layer, Conv2D):
            lr_mult[layer.name + '/kernel:0'] = 0.001
            lr_mult[layer.name + '/bias:0'] = 0.001

    # Set learning rate to 10 for the last two dense layers
    for layer in model_VGG19.layers[23:]:
        lr_mult[layer.name + '/kernel:0'] = 0.001
        lr_mult[layer.name + '/bias:0'] = 0.001

    # Use SGD optimizer with initial learning rate of 1e-4
    optimizer = SGD(learning_rate=1e-4, momentum=0.9)

    # Compile the model
```

```

    model_VGG19.compile(optimizer=optimizer, loss='categorical_crossentropy',
↪metrics=['accuracy'], loss_weights=[1.] + [10.]*2)

    return model_VGG19

accuracy_values = []
precision_values = []
recall_values = []
f1_values = []

train_acc_list = []
val_acc_list = []
train_loss_list = []
val_loss_list = []

# KFold Cross-validation
n_splits = 5
kf = KFold(n_splits=n_splits, shuffle=True, random_state=777)
fold_var = 1

for train_index, test_index in kf.split(img_data):
    X_train, X_test = img_data[train_index], img_data[test_index]
    y_train, y_test = labels_categorical[train_index],
↪labels_categorical[test_index]

    model_VGG19 = build_model()

    #early_stop = EarlyStopping(monitor='val_loss', patience=3, verbose=1,
↪mode='min')
    save_path='/content/drive/MyDrive/Deep_learning_projects/SheepFull_Data/
↪SheepFace_three.keras'
    checkpoint =
↪ModelCheckpoint(save_path,monitor='val_loss',save_best_only=True,verbose=0)
    history = model_VGG19.fit(X_train, y_train, batch_size=10, epochs=10,
↪verbose=1, validation_split=0.2, callbacks=[checkpoint])

    print("\n")
    print(f'Evaluating the Test metrics')
    model_VGG19=load_model(save_path)
    # Evaluate the model
    scores = model_VGG19.evaluate(X_test, y_test, verbose=1)
    print(f'Test loss for fold {fold_var}: {scores[0]}')
    print(f'Test accuracy for fold {fold_var}: {scores[1]}')

    accuracy_values.append(scores[1]) # Appending accuracy to the list

```

```

# Get predictions from the model
y_true = np.argmax(y_test, axis=1) # Convert one-hot encoded labels back
↳to categorical labels
y_pred = np.argmax(model_VGG19.predict(X_test), axis=1) # Get predictions
↳from the model

# Calculate and append evaluation metrics for this fold
precision = precision_score(y_true, y_pred, average='weighted',
↳zero_division=1)
recall = recall_score(y_true, y_pred, average='weighted')
f1 = f1_score(y_true, y_pred, average='weighted')

precision_values.append(precision)
recall_values.append(recall)
f1_values.append(f1)

# Print evaluation metrics for this fold
print(f'Precision for fold {fold_var}: {precision}')
print(f'Recall for fold {fold_var}: {recall}')
print(f'F1 Score for fold {fold_var}: {f1}')
print("\n")

train_acc_list.append(history.history['accuracy'])
val_acc_list.append(history.history['val_accuracy'])
train_loss_list.append(history.history['loss'])
val_loss_list.append(history.history['val_loss'])

# Increment the fold number
fold_var += 1

# Calculate standard deviation of accuracy across folds
accuracy_std_dev = np.std(accuracy_values)
print(f'Standard deviation of accuracy across folds: {accuracy_std_dev}')

# Calculate average of accuracy, precision, recall, and F1 across folds
avg_accuracy = np.mean(accuracy_values)
avg_precision = np.mean(precision_values)
avg_recall = np.mean(recall_values)
avg_f1 = np.mean(f1_values)

print(f'Average Accuracy across folds: {avg_accuracy}')
print(f'Average Precision across folds: {avg_precision}')
print(f'Average Recall across folds: {avg_recall}')
print(f'Average F1 Score across folds: {avg_f1}')

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5)

```

80134624/80134624 [=====] - 4s 0us/step
Epoch 1/10
104/104 [=====] - 22s 154ms/step - loss: 1.6248 -
accuracy: 0.5830 - val_loss: 0.6109 - val_accuracy: 0.7683
Epoch 2/10
104/104 [=====] - 11s 104ms/step - loss: 0.4059 -
accuracy: 0.8542 - val_loss: 0.3482 - val_accuracy: 0.8571
Epoch 3/10
104/104 [=====] - 10s 95ms/step - loss: 0.1213 -
accuracy: 0.9614 - val_loss: 0.4795 - val_accuracy: 0.8263
Epoch 4/10
104/104 [=====] - 11s 106ms/step - loss: 0.0400 -
accuracy: 0.9923 - val_loss: 0.2234 - val_accuracy: 0.9266
Epoch 5/10
104/104 [=====] - 10s 97ms/step - loss: 0.0071 -
accuracy: 1.0000 - val_loss: 0.2361 - val_accuracy: 0.9305
Epoch 6/10
104/104 [=====] - 10s 97ms/step - loss: 0.0034 -
accuracy: 1.0000 - val_loss: 0.2241 - val_accuracy: 0.9266
Epoch 7/10
104/104 [=====] - 11s 107ms/step - loss: 0.0021 -
accuracy: 1.0000 - val_loss: 0.2182 - val_accuracy: 0.9266
Epoch 8/10
104/104 [=====] - 10s 98ms/step - loss: 0.0017 -
accuracy: 1.0000 - val_loss: 0.2192 - val_accuracy: 0.9305
Epoch 9/10
104/104 [=====] - 10s 98ms/step - loss: 0.0014 -
accuracy: 1.0000 - val_loss: 0.2218 - val_accuracy: 0.9344
Epoch 10/10
104/104 [=====] - 11s 107ms/step - loss: 0.0012 -
accuracy: 1.0000 - val_loss: 0.2165 - val_accuracy: 0.9382

```

Evaluating the Test metrics

```

11/11 [=====] - 10s 291ms/step - loss: 0.2751 -
accuracy: 0.9074
Test loss for fold 1: 0.27506136894226074
Test accuracy for fold 1: 0.9074074029922485
11/11 [=====] - 2s 143ms/step
Precision for fold 1: 0.9083414893583627
Recall for fold 1: 0.9074074074074074
F1 Score for fold 1: 0.9076904694338656

```

```

Epoch 1/10
104/104 [=====] - 13s 113ms/step - loss: 1.6889 -
accuracy: 0.5618 - val_loss: 0.7950 - val_accuracy: 0.6988
Epoch 2/10

```



104/104 [=====] - 11s 108ms/step - loss: 0.5662 -  
accuracy: 0.7867 - val\_loss: 0.4806 - val\_accuracy: 0.8340  
Epoch 3/10  
104/104 [=====] - 11s 108ms/step - loss: 0.2349 -  
accuracy: 0.9112 - val\_loss: 0.2357 - val\_accuracy: 0.9189  
Epoch 4/10  
104/104 [=====] - 11s 107ms/step - loss: 0.0917 -  
accuracy: 0.9759 - val\_loss: 0.2065 - val\_accuracy: 0.9266  
Epoch 5/10  
104/104 [=====] - 11s 107ms/step - loss: 0.0405 -  
accuracy: 0.9913 - val\_loss: 0.1976 - val\_accuracy: 0.9344  
Epoch 6/10  
104/104 [=====] - 11s 106ms/step - loss: 0.0143 -  
accuracy: 0.9981 - val\_loss: 0.1843 - val\_accuracy: 0.9459  
Epoch 7/10  
104/104 [=====] - 11s 107ms/step - loss: 0.0077 -  
accuracy: 1.0000 - val\_loss: 0.1702 - val\_accuracy: 0.9537  
Epoch 8/10  
104/104 [=====] - 11s 106ms/step - loss: 0.0036 -  
accuracy: 1.0000 - val\_loss: 0.1670 - val\_accuracy: 0.9498  
Epoch 9/10  
104/104 [=====] - 10s 97ms/step - loss: 0.0024 -  
accuracy: 1.0000 - val\_loss: 0.1683 - val\_accuracy: 0.9537  
Epoch 10/10  
104/104 [=====] - 11s 107ms/step - loss: 0.0018 -  
accuracy: 1.0000 - val\_loss: 0.1669 - val\_accuracy: 0.9459

#### Evaluating the Test metrics

11/11 [=====] - 2s 134ms/step - loss: 0.2931 -  
accuracy: 0.9352  
Test loss for fold 2: 0.29307472705841064  
Test accuracy for fold 2: 0.9351851940155029  
11/11 [=====] - 2s 141ms/step  
Precision for fold 2: 0.9375098284515365  
Recall for fold 2: 0.9351851851851852  
F1 Score for fold 2: 0.9348004574033981

Epoch 1/10  
104/104 [=====] - 13s 113ms/step - loss: 1.8972 -  
accuracy: 0.5203 - val\_loss: 0.8484 - val\_accuracy: 0.6371  
Epoch 2/10  
104/104 [=====] - 11s 108ms/step - loss: 0.5761 -  
accuracy: 0.7867 - val\_loss: 0.3972 - val\_accuracy: 0.8571  
Epoch 3/10  
104/104 [=====] - 10s 98ms/step - loss: 0.3207 -  
accuracy: 0.8851 - val\_loss: 0.4739 - val\_accuracy: 0.8456

Epoch 4/10  
104/104 [=====] - 11s 107ms/step - loss: 0.1219 -  
accuracy: 0.9614 - val\_loss: 0.2305 - val\_accuracy: 0.9266  
Epoch 5/10  
104/104 [=====] - 10s 97ms/step - loss: 0.0487 -  
accuracy: 0.9865 - val\_loss: 0.2321 - val\_accuracy: 0.9189  
Epoch 6/10  
104/104 [=====] - 11s 107ms/step - loss: 0.0185 -  
accuracy: 0.9971 - val\_loss: 0.1695 - val\_accuracy: 0.9382  
Epoch 7/10  
104/104 [=====] - 11s 107ms/step - loss: 0.0082 -  
accuracy: 0.9990 - val\_loss: 0.1667 - val\_accuracy: 0.9421  
Epoch 8/10  
104/104 [=====] - 10s 97ms/step - loss: 0.0034 -  
accuracy: 1.0000 - val\_loss: 0.1907 - val\_accuracy: 0.9344  
Epoch 9/10  
104/104 [=====] - 10s 97ms/step - loss: 0.0021 -  
accuracy: 1.0000 - val\_loss: 0.1859 - val\_accuracy: 0.9421  
Epoch 10/10  
104/104 [=====] - 10s 97ms/step - loss: 0.0017 -  
accuracy: 1.0000 - val\_loss: 0.1821 - val\_accuracy: 0.9382

Evaluating the Test metrics

11/11 [=====] - 2s 135ms/step - loss: 0.2426 -  
accuracy: 0.9198  
Test loss for fold 3: 0.24258004128932953  
Test accuracy for fold 3: 0.9197530746459961  
11/11 [=====] - 2s 141ms/step  
Precision for fold 3: 0.9225461497023901  
Recall for fold 3: 0.9197530864197531  
F1 Score for fold 3: 0.9202892779027154

Epoch 1/10  
104/104 [=====] - 13s 113ms/step - loss: 1.3616 -  
accuracy: 0.6390 - val\_loss: 0.5852 - val\_accuracy: 0.8031  
Epoch 2/10  
104/104 [=====] - 11s 108ms/step - loss: 0.2675 -  
accuracy: 0.9093 - val\_loss: 0.2763 - val\_accuracy: 0.9073  
Epoch 3/10  
104/104 [=====] - 11s 108ms/step - loss: 0.0658 -  
accuracy: 0.9875 - val\_loss: 0.2009 - val\_accuracy: 0.9228  
Epoch 4/10  
104/104 [=====] - 11s 108ms/step - loss: 0.0186 -  
accuracy: 0.9971 - val\_loss: 0.1904 - val\_accuracy: 0.9305  
Epoch 5/10  
104/104 [=====] - 11s 107ms/step - loss: 0.0063 -

accuracy: 1.0000 - val\_loss: 0.1654 - val\_accuracy: 0.9575  
Epoch 6/10  
104/104 [=====] - 11s 107ms/step - loss: 0.0030 -  
accuracy: 1.0000 - val\_loss: 0.1562 - val\_accuracy: 0.9537  
Epoch 7/10  
104/104 [=====] - 10s 97ms/step - loss: 0.0020 -  
accuracy: 1.0000 - val\_loss: 0.1567 - val\_accuracy: 0.9575  
Epoch 8/10  
104/104 [=====] - 11s 107ms/step - loss: 0.0016 -  
accuracy: 1.0000 - val\_loss: 0.1546 - val\_accuracy: 0.9575  
Epoch 9/10  
104/104 [=====] - 10s 97ms/step - loss: 0.0013 -  
accuracy: 1.0000 - val\_loss: 0.1560 - val\_accuracy: 0.9575  
Epoch 10/10  
104/104 [=====] - 10s 97ms/step - loss: 0.0012 -  
accuracy: 1.0000 - val\_loss: 0.1551 - val\_accuracy: 0.9575

Evaluating the Test metrics

11/11 [=====] - 2s 134ms/step - loss: 0.1838 -  
accuracy: 0.9198  
Test loss for fold 4: 0.18379390239715576  
Test accuracy for fold 4: 0.9197530746459961  
11/11 [=====] - 2s 142ms/step  
Precision for fold 4: 0.9203762688044644  
Recall for fold 4: 0.9197530864197531  
F1 Score for fold 4: 0.91991264595853

Epoch 1/10  
104/104 [=====] - 13s 113ms/step - loss: 1.4446 -  
accuracy: 0.6207 - val\_loss: 0.5601 - val\_accuracy: 0.7615  
Epoch 2/10  
104/104 [=====] - 11s 108ms/step - loss: 0.3344 -  
accuracy: 0.8755 - val\_loss: 0.3099 - val\_accuracy: 0.9038  
Epoch 3/10  
104/104 [=====] - 11s 108ms/step - loss: 0.0859 -  
accuracy: 0.9730 - val\_loss: 0.2175 - val\_accuracy: 0.9192  
Epoch 4/10  
104/104 [=====] - 11s 108ms/step - loss: 0.0155 -  
accuracy: 0.9981 - val\_loss: 0.1613 - val\_accuracy: 0.9462  
Epoch 5/10  
104/104 [=====] - 11s 107ms/step - loss: 0.0033 -  
accuracy: 1.0000 - val\_loss: 0.1348 - val\_accuracy: 0.9577  
Epoch 6/10  
104/104 [=====] - 11s 108ms/step - loss: 0.0019 -  
accuracy: 1.0000 - val\_loss: 0.1344 - val\_accuracy: 0.9577  
Epoch 7/10

```

104/104 [=====] - 11s 107ms/step - loss: 0.0015 -
accuracy: 1.0000 - val_loss: 0.1339 - val_accuracy: 0.9500
Epoch 8/10
104/104 [=====] - 12s 112ms/step - loss: 0.0012 -
accuracy: 1.0000 - val_loss: 0.1313 - val_accuracy: 0.9577
Epoch 9/10
104/104 [=====] - 11s 107ms/step - loss: 0.0010 -
accuracy: 1.0000 - val_loss: 0.1310 - val_accuracy: 0.9577
Epoch 10/10
104/104 [=====] - 10s 97ms/step - loss: 8.8104e-04 -
accuracy: 1.0000 - val_loss: 0.1314 - val_accuracy: 0.9538

```

Evaluating the Test metrics

```

11/11 [=====] - 3s 266ms/step - loss: 0.1983 -
accuracy: 0.9226
Test loss for fold 5: 0.19832882285118103
Test accuracy for fold 5: 0.9226006269454956
11/11 [=====] - 2s 142ms/step
Precision for fold 5: 0.9223827986465892
Recall for fold 5: 0.9226006191950464
F1 Score for fold 5: 0.92243921584306

```

```

Standard deviation of accuracy across folds: 0.008850220332613429
Average Accuracy across folds: 0.9209398746490478
Average Precision across folds: 0.9222313069926686
Average Recall across folds: 0.9209398769254291
Average F1 Score across folds: 0.9210264133083138

```

```

[9]: import matplotlib.pyplot as plt

# Determine the maximum length among all lists
max_length = max(len(train_acc) for train_acc in train_acc_list)

# Pad the shorter lists with zeros to match the maximum length
for i in range(n_splits):
    train_acc_list[i] += [0] * (max_length - len(train_acc_list[i]))
    val_acc_list[i] += [0] * (max_length - len(val_acc_list[i]))
    train_loss_list[i] += [0] * (max_length - len(train_loss_list[i]))
    val_loss_list[i] += [0] * (max_length - len(val_loss_list[i]))

# Plotting
epochs = range(1, max_length + 1)

plt.figure(figsize=(12, 6))

```

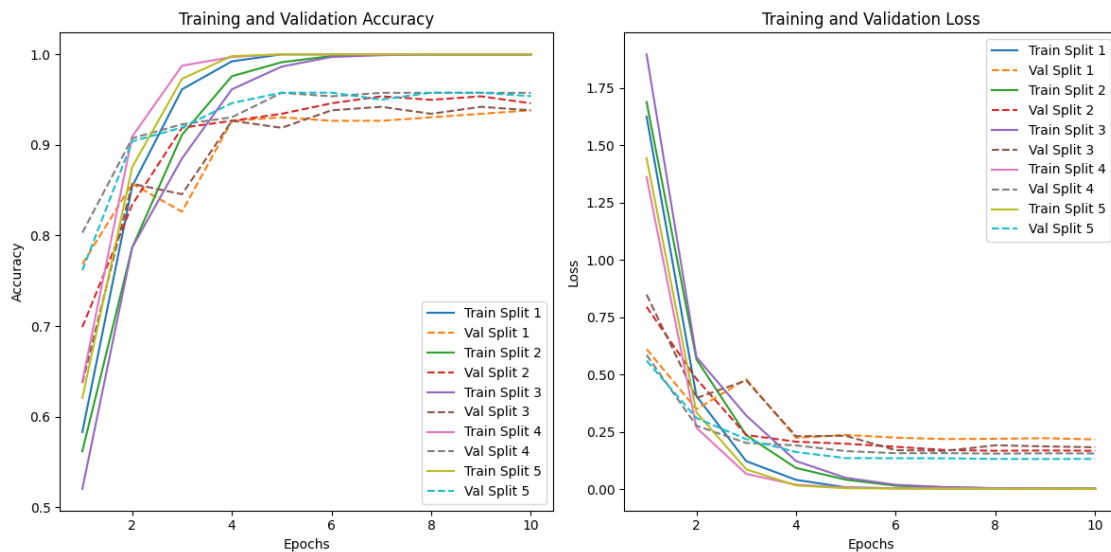
```

# Plotting Training and Validation Accuracy
plt.subplot(1, 2, 1)
for i in range(n_splits):
    plt.plot(epochs, train_acc_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_acc_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Plotting Training and Validation Loss
plt.subplot(1, 2, 2)
for i in range(n_splits):
    plt.plot(epochs, train_loss_list[i], label=f'Train Split {i+1}')
    plt.plot(epochs, val_loss_list[i], label=f'Val Split {i+1}', linestyle='--')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()

```



```

[19]: from keras.applications.vgg16 import VGG16, preprocess_input
      from sklearn.svm import SVC
      from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import StandardScaler
      from sklearn.preprocessing import LabelEncoder

```

```

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import StratifiedKFold
import numpy as np

encoder = LabelEncoder()
labels_encoded = encoder.fit_transform(labels)

# Load the pre-trained VGG16 model without the top (classification) layers
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze the layers of the pre-trained model
for layer in base_model.layers:
    layer.trainable = False

# Define SVM pipeline
svm_pipeline = Pipeline([('scaler', StandardScaler()), ('svm', SVC(kernel='linear', probability=True))])

# Initialize lists to store evaluation metrics
accuracy_values = []
precision_values = []
recall_values = []
f1_values = []

# Perform k-fold cross-validation
n_splits = 5
skf = StratifiedKFold(n_splits=n_splits, shuffle=True, random_state=777)
fold_var = 1

for train_index, test_index in skf.split(img_data, labels_encoded): # Ensure labels_encoded is defined correctly as the non-categorical labels
    X_train, X_test = img_data[train_index], img_data[test_index]
    y_train, y_test = labels_encoded[train_index], labels_encoded[test_index]

    # Extract features using the VGG16 model and flatten them
    X_train_features = base_model.predict(preprocess_input(X_train))
    X_train_features = X_train_features.reshape(X_train_features.shape[0], -1)

    X_test_features = base_model.predict(preprocess_input(X_test))
    X_test_features = X_test_features.reshape(X_test_features.shape[0], -1)

    # Fit and transform the training data with StandardScaler and train the SVM
    svm_pipeline.fit(X_train_features, y_train)

    # Predict using the trained SVM pipeline

```

```

predictions = svm_pipeline.predict(X_test_features)
print(f'Fold {fold_var}:')
print(f'Accuracy: {accuracy_score(y_test, predictions)}')
print(f'Precision: {precision_score(y_test, predictions,
↪average="weighted")}')
print(f'Recall: {recall_score(y_test, predictions, average="weighted")}')
print(f'F1 Score: {f1_score(y_test, predictions, average="weighted")}')

# Calculate evaluation metrics
accuracy_values.append(accuracy_score(y_test, predictions))
precision_values.append(precision_score(y_test, predictions,
↪average='weighted'))
recall_values.append(recall_score(y_test, predictions, average='weighted'))
f1_values.append(f1_score(y_test, predictions, average='weighted'))

# Increment the fold number
fold_var += 1

# Calculate average metrics across all folds
avg_accuracy = np.mean(accuracy_values)
avg_precision = np.mean(precision_values)
avg_recall = np.mean(recall_values)
avg_f1 = np.mean(f1_values)

print(f'Average test Accuracy across folds: {avg_accuracy}')
print(f'Average Precision across folds: {avg_precision}')
print(f'Average Recall across folds: {avg_recall}')
print(f'Average F1 Score across folds: {avg_f1}')

```

41/41 [=====] - 10s 250ms/step

11/11 [=====] - 2s 148ms/step

Fold 1:

Accuracy: 0.9598765432098766

Precision: 0.9620569043250449

Recall: 0.9598765432098766

F1 Score: 0.9595809528875071

41/41 [=====] - 6s 138ms/step

11/11 [=====] - 1s 141ms/step

Fold 2:

Accuracy: 0.9290123456790124

Precision: 0.9307860902476693

Recall: 0.9290123456790124

F1 Score: 0.9292381459241569

41/41 [=====] - 6s 144ms/step

11/11 [=====] - 1s 143ms/step

Fold 3:

Accuracy: 0.9660493827160493

```

Precision: 0.9663775683757447
Recall: 0.9660493827160493
F1 Score: 0.9658663233744746
41/41 [=====] - 6s 140ms/step
11/11 [=====] - 1s 142ms/step
Fold 4:
Accuracy: 0.9598765432098766
Precision: 0.9613916475549054
Recall: 0.9598765432098766
F1 Score: 0.9598198652756754
41/41 [=====] - 10s 250ms/step
11/11 [=====] - 1s 144ms/step
Fold 5:
Accuracy: 0.9597523219814241
Precision: 0.9601101715746051
Recall: 0.9597523219814241
F1 Score: 0.9598046349590906
Average test Accuracy across folds: 0.9549134273592479
Average Precision across folds: 0.956144476415594
Average Recall across folds: 0.9549134273592479
Average F1 Score across folds: 0.9548619844841809

```

```

[5]: !jupyter nbconvert --to pdf '/content/drive/MyDrive/Colab_Notebooks/
     ↪Sheep_Full_Images_Classification.ipynb'

```

```

[NbConvertApp] Converting notebook
/content/drive/MyDrive/Colab_Notebooks/Sheep_Full_Images_Classification.ipynb to
pdf
[NbConvertApp] Support files will be in Sheep_Full_Images_Classification_files/
[NbConvertApp] Making directory ./Sheep_Full_Images_Classification_files
[NbConvertApp] Making directory ./Sheep_Full_Images_Classification_files
[NbConvertApp] Making directory ./Sheep_Full_Images_Classification_files
[NbConvertApp] Writing 178570 bytes to notebook.tex
[NbConvertApp] Building PDF
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no
citations
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 544764 bytes to
/content/drive/MyDrive/Colab_Notebooks/Sheep_Full_Images_Classification.pdf

```